

# CSC 226

## Algorithms and Data Structures: II More Red-Black Trees

Tianming Wei

[twei@uvic.ca](mailto:twei@uvic.ca)

ECS 466

# Red-black trees and 2-3 trees

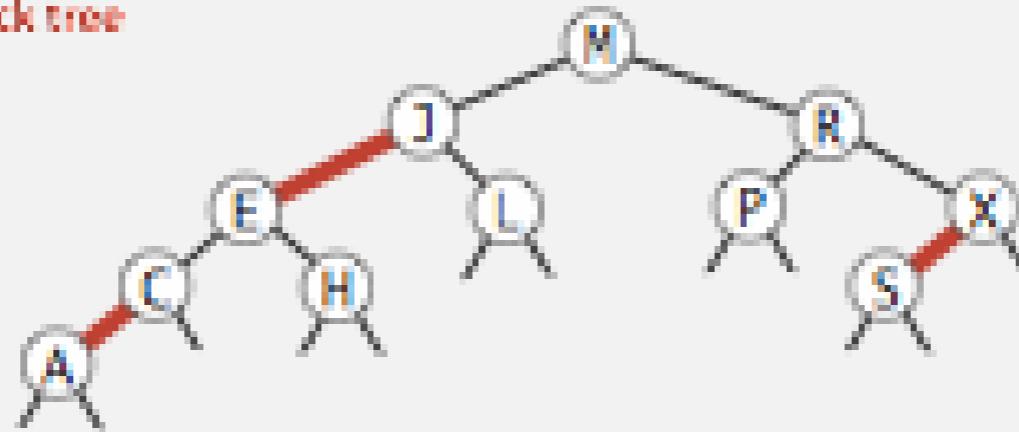
- There is a 1-1 correspondence between red-black BSTs and 2-3 trees.
- To see this, imagine that red links are collapsed: the collapsed nodes correspond to 3-nodes, all others to 2-nodes.

# Left-leaning red-black BSTs: 1-1 correspondence with 2-3 trees

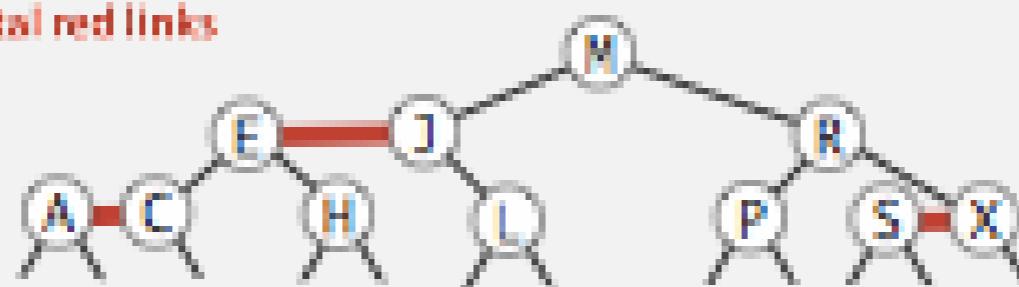
---

**Key property.** 1-1 correspondence between 2-3 and LLRB.

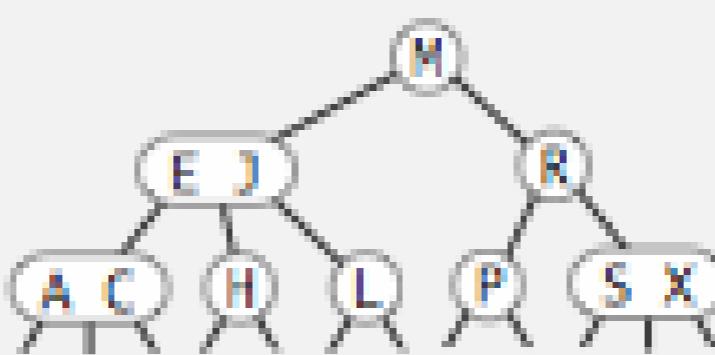
red-black tree



horizontal red links



2-3 tree



# We show: For every red-black tree there is a 2-3 tree

- Given a red-black tree  $T$ , we build a 2-3 tree  $T'$
- The nodes of the 2-3 tree  $T'$  are obtained as follows.
  - for every node  $v_{rb}$  in the red-black tree with key  $k$  that is incident to black edges only create a 2-node  $v_{23}$  for the 2-3 tree with key  $k$ . That is:  $23(v_{rb}) = v_{23}$
  - for every **red** edge and its incident two nodes  $v_{rb}$  and  $w_{rb}$ —where  $w_{rb}$  is the parent of  $v_{rb}$ —containing keys  $k_1$  and  $k_2$ , respectively, create a 3-node  $vw_{23}$  containing keys  $k_1$  and  $k_2$  (with  $k_1$  being the left entry). That is:  $23(v_{rb}) = 23(w_{rb}) = vw_{23}$

# We show: For every red-black tree there is a 2-3 tree

- The edges of the 2-3 tree  $T'$  are obtained as follows
  - Let  $u_{rb}v_{rb}$  be a **black** edge in the red-black tree. Then  $23(u_{rb})23(v_{rb})$  is an edge in the 2-3 tree. Further,
    - if  $u_{rb}$  is the left child of  $v_{rb}$  then  $23(u_{rb})$  is the left child of  $23(v_{rb})$
    - else if  $u_{rb}$  is the right child of  $v_{rb}$ , and the edge between  $v_{rb}$  and its parent is red, then  $23(u_{rb})$  is  $23(v_{rb})$ 's middle child
    - else  $23(u_{rb})$  is the right child of  $23(v_{rb})$

# We show: For every 2-3 tree there is a red-black tree

- Given a 2-3 tree  $T$ , we build a red-black tree  $T'$
- The nodes of the red-black tree  $T'$  are obtained as follows.
  - for every 2-node  $v_{23}$  in the red-black tree with key  $k$  create node  $v_{rb}$  with key  $k$ .
  - for every 3-node  $vw_{23}$  with keys  $k_1$  and  $k_2$ ,  $k_1 \leq k_2$ , create nodes  $v_{rb}$  and  $w_{rb}$  containing keys  $k_1$  and  $k_2$ , respectively.

# We show: For every red-black tree there is a 2-3 tree

- The edges of the red-black tree are obtained as follows
  - for every remaining edge between 2-node  $w_{23}$  and child  $v_{23}$  the corresponding red-black tree nodes are connected by a black edge
  - for every 3-node  $vw_{23}$  with keys  $k_1$  and  $k_2$ ,  $k_1 \leq k_2$ , and its corresponding red-black tree nodes  $v_{rb}$  with key  $k_1$  and  $w_{rb}$  with key  $k_2$  we define recursively:
    - add **red** edge  $v_{rb}w_{rb}$  in red-black tree  $T'$  with  $v_{rb}$  being  $w_{rb}$ 's left child
    - the red-black tree of  $vw_{23}$ 's left subtree is  $v_{rb}$ 's left subtree, the red-black tree of  $vw_{23}$ 's middle subtree is  $v_{rb}$ 's right subtree, and the red-black tree of  $vw_{23}$ 's right subtree is  $w_{rb}$ 's right subtree

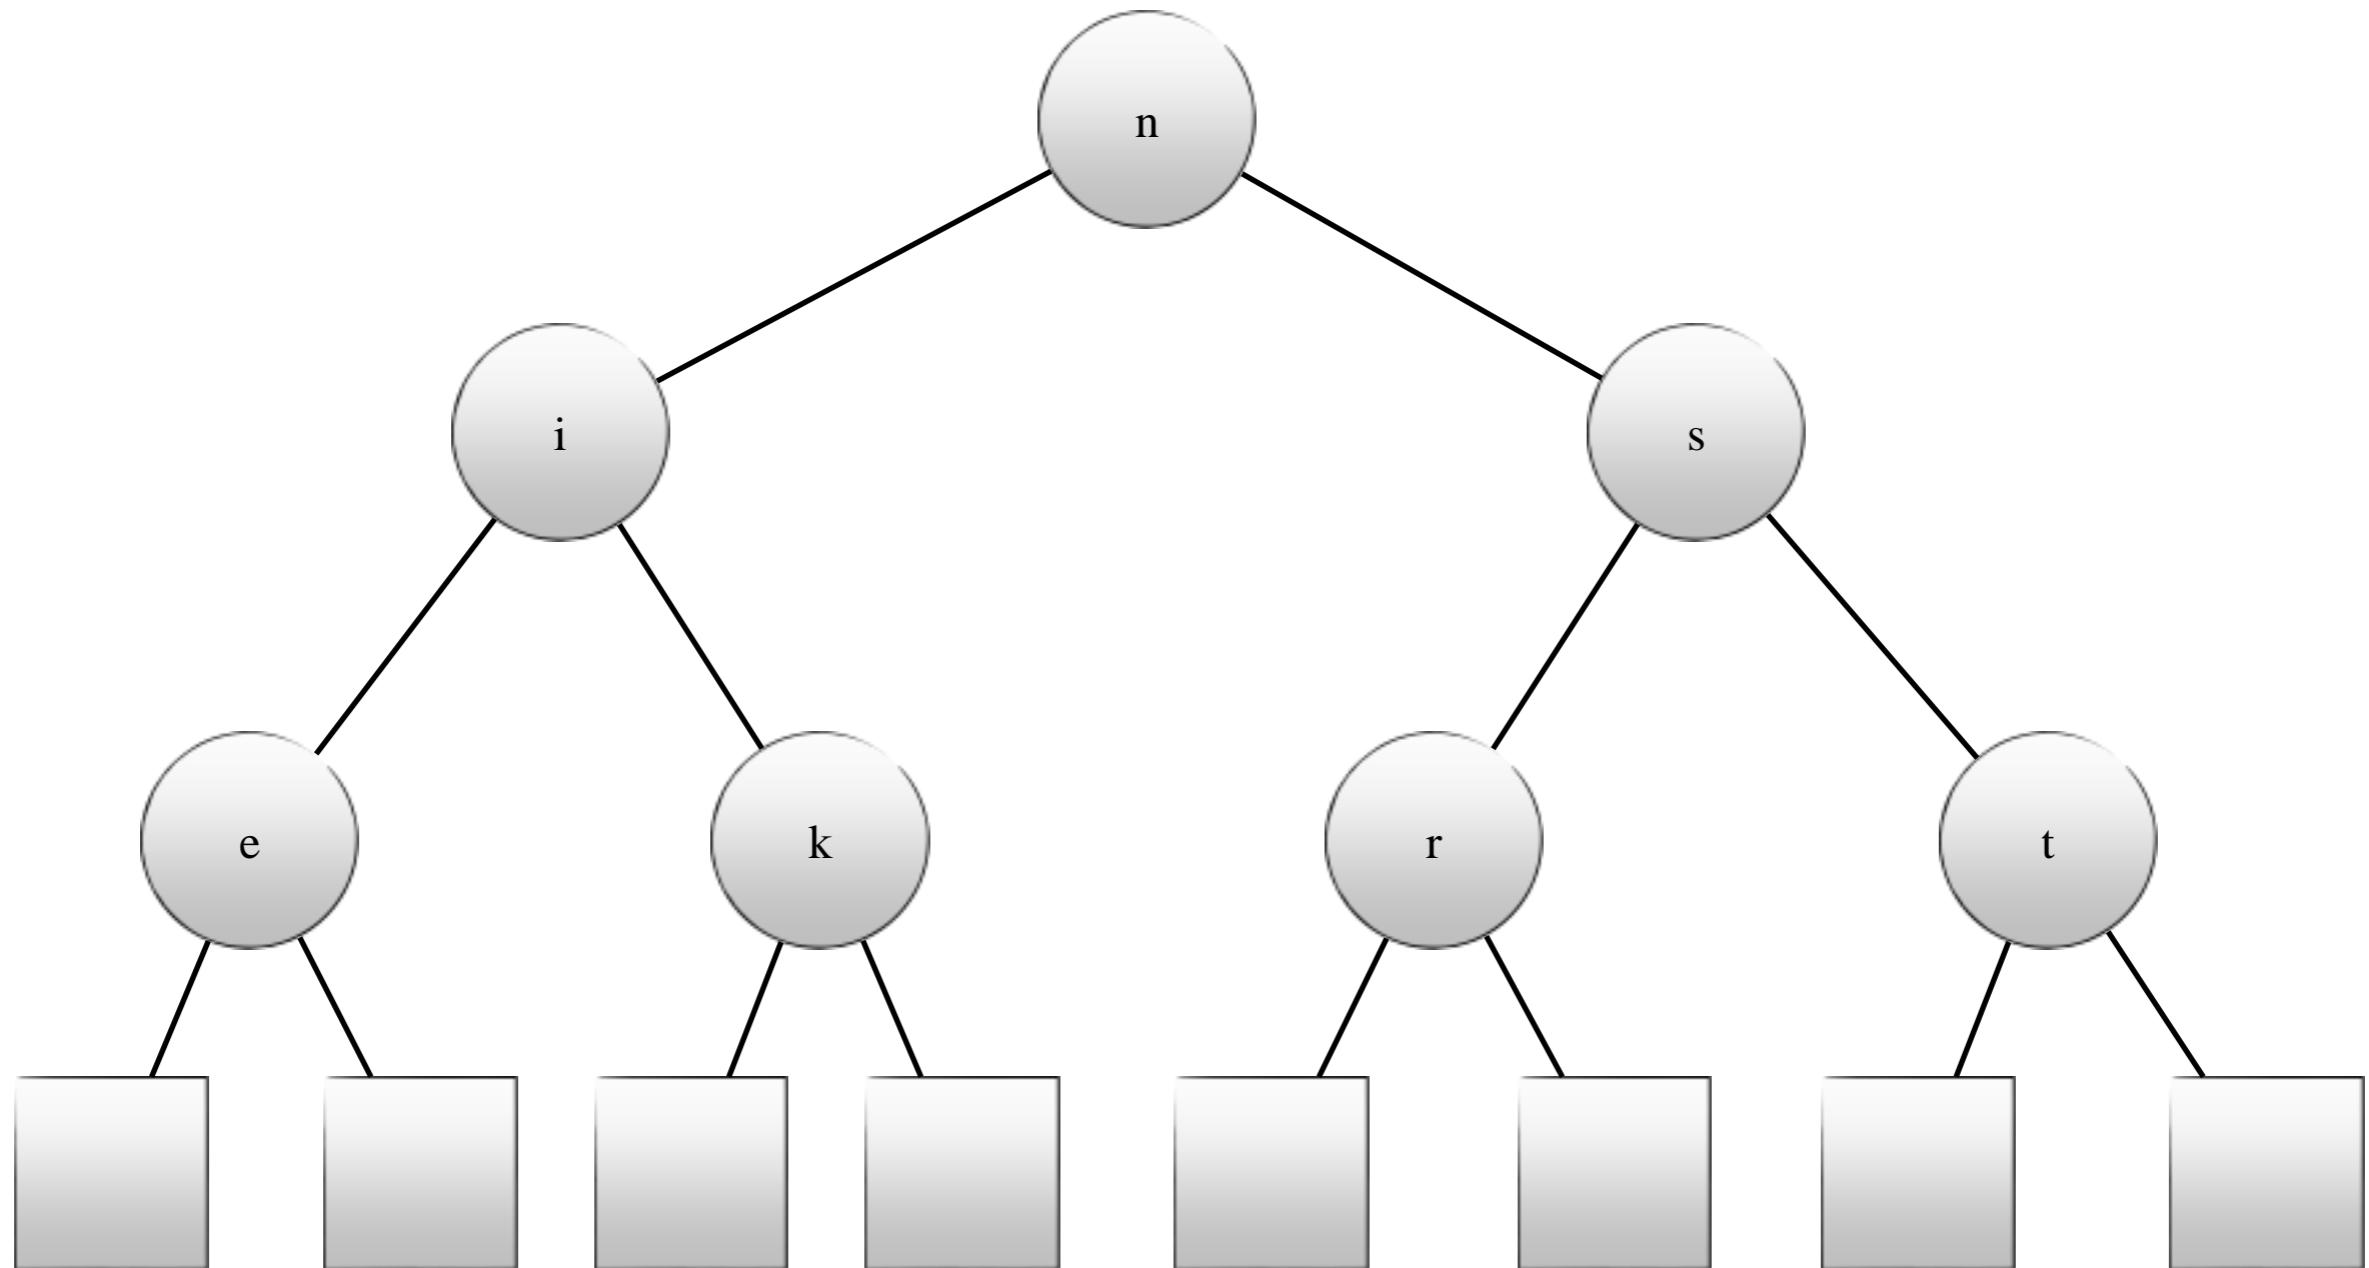
# The height of a red-black tree with $n$ keys is $O(\log(n))$

- *Proof Idea.* The (black) path length from leaf to root is the same as the height of its corresponding 2-3 tree, that is  $O(\log(n))$ . The height of a red-black tree can be twice the height of the 2-3 tree, namely in the case that every second edge in a path from leaf to root is red. The height remains  $O(\log(n))$ .

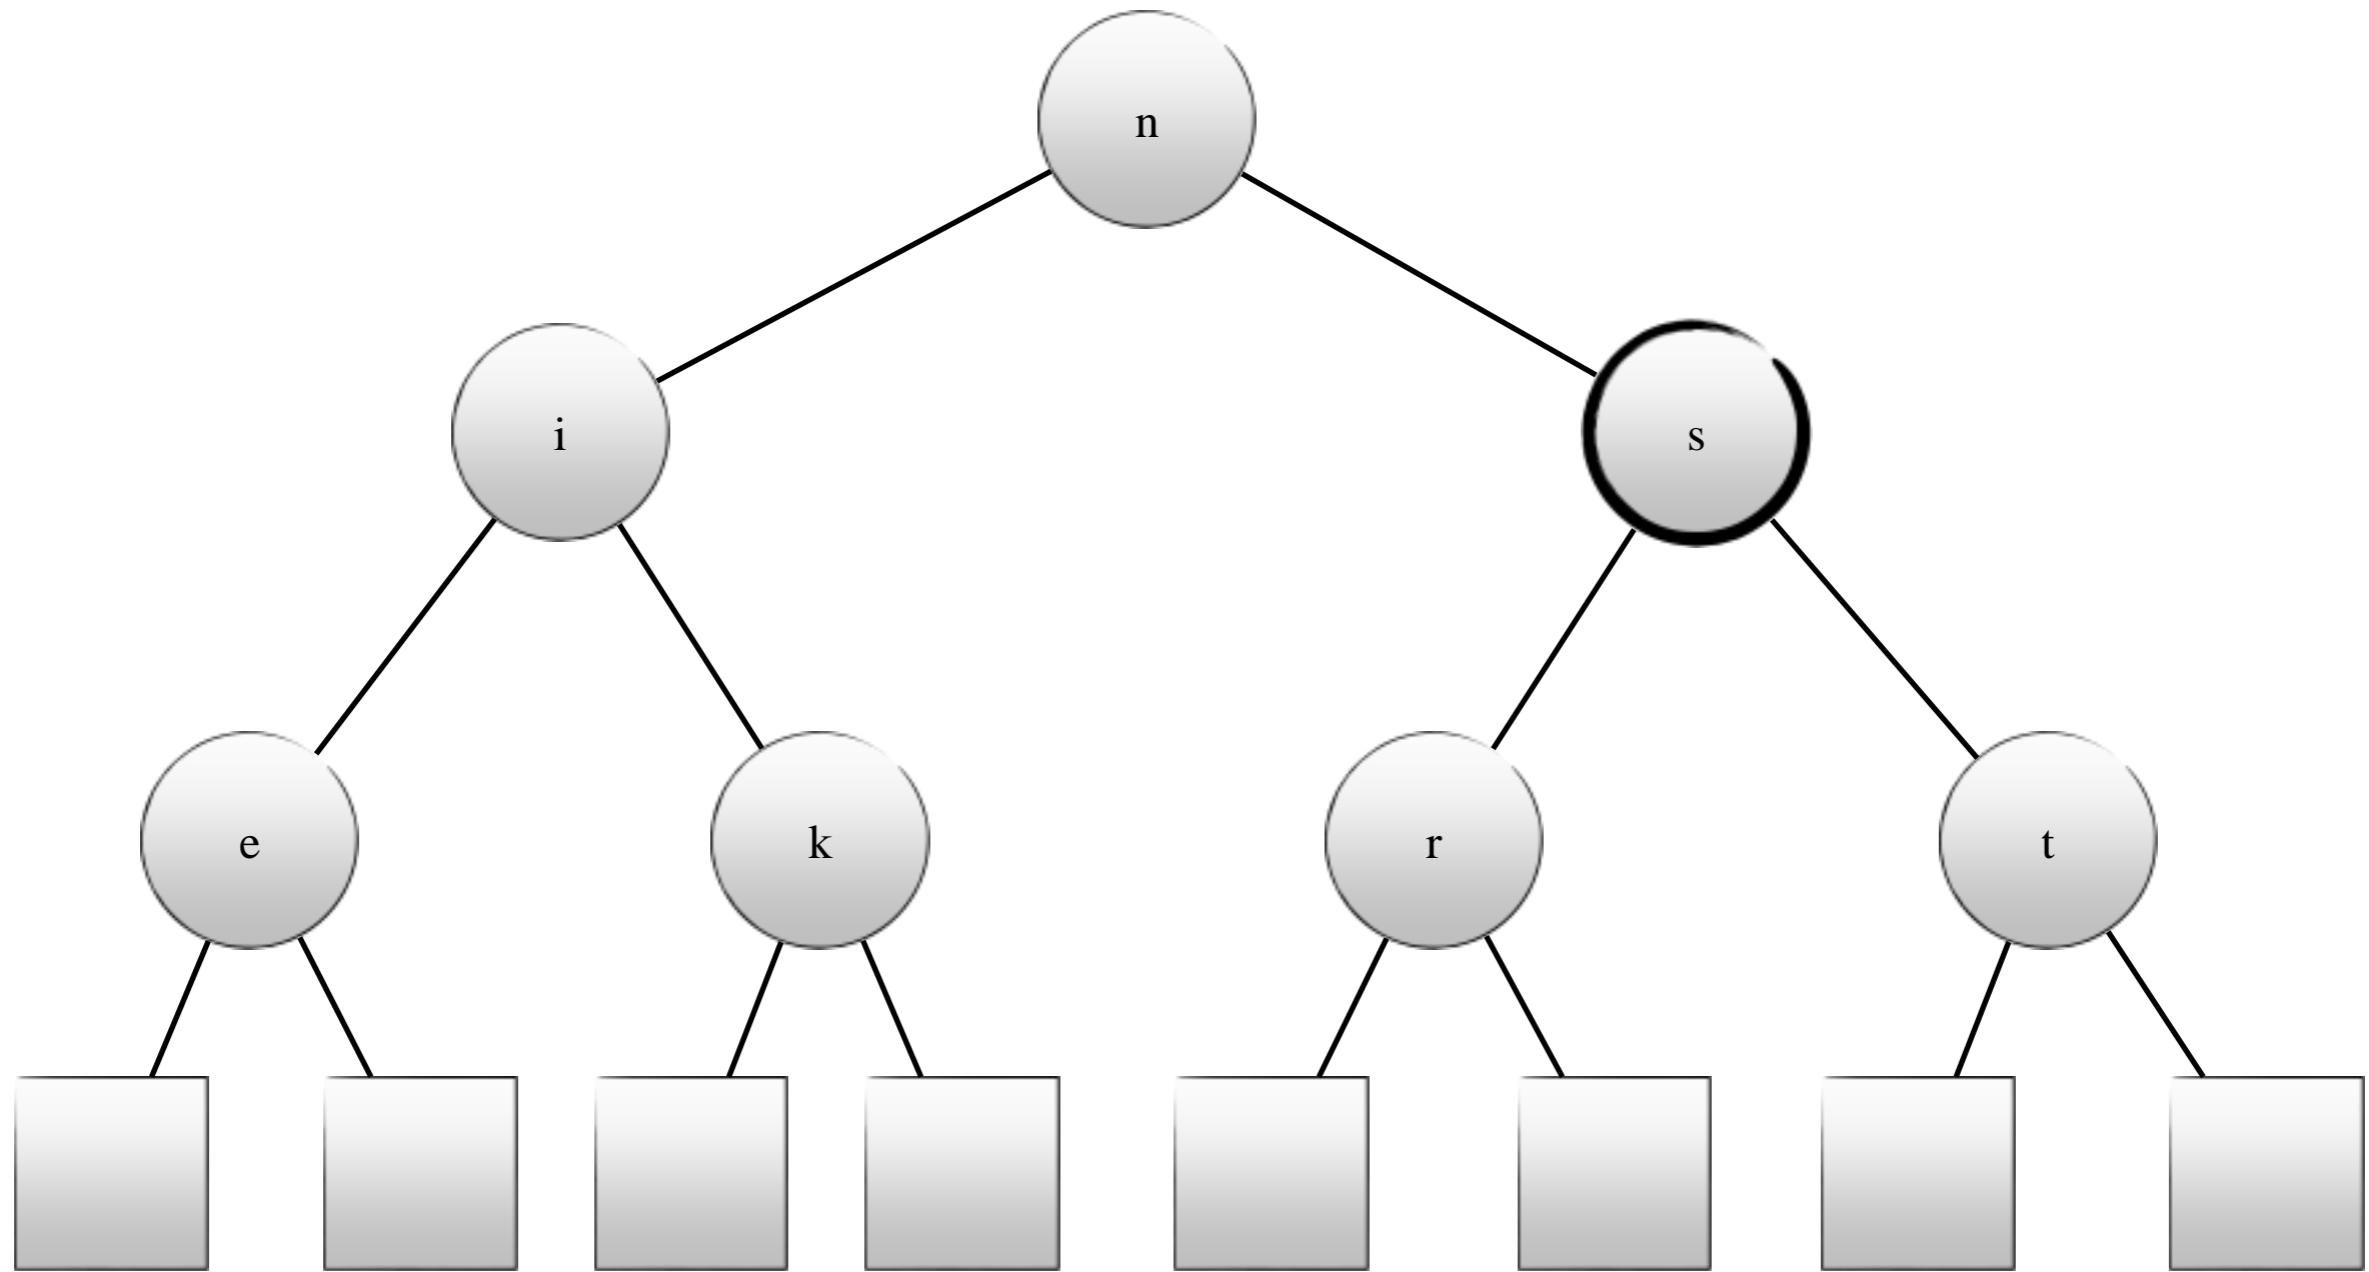
# Deletion from 2-3/red-black trees

- Deletion is more problematic than insertion. Why?
- During insertion we can only cause problems with the colouring of the nodes
  - i.e. violate the left red and/or one red edge properties.
- During deletion we can cause black-depth imbalance as well.

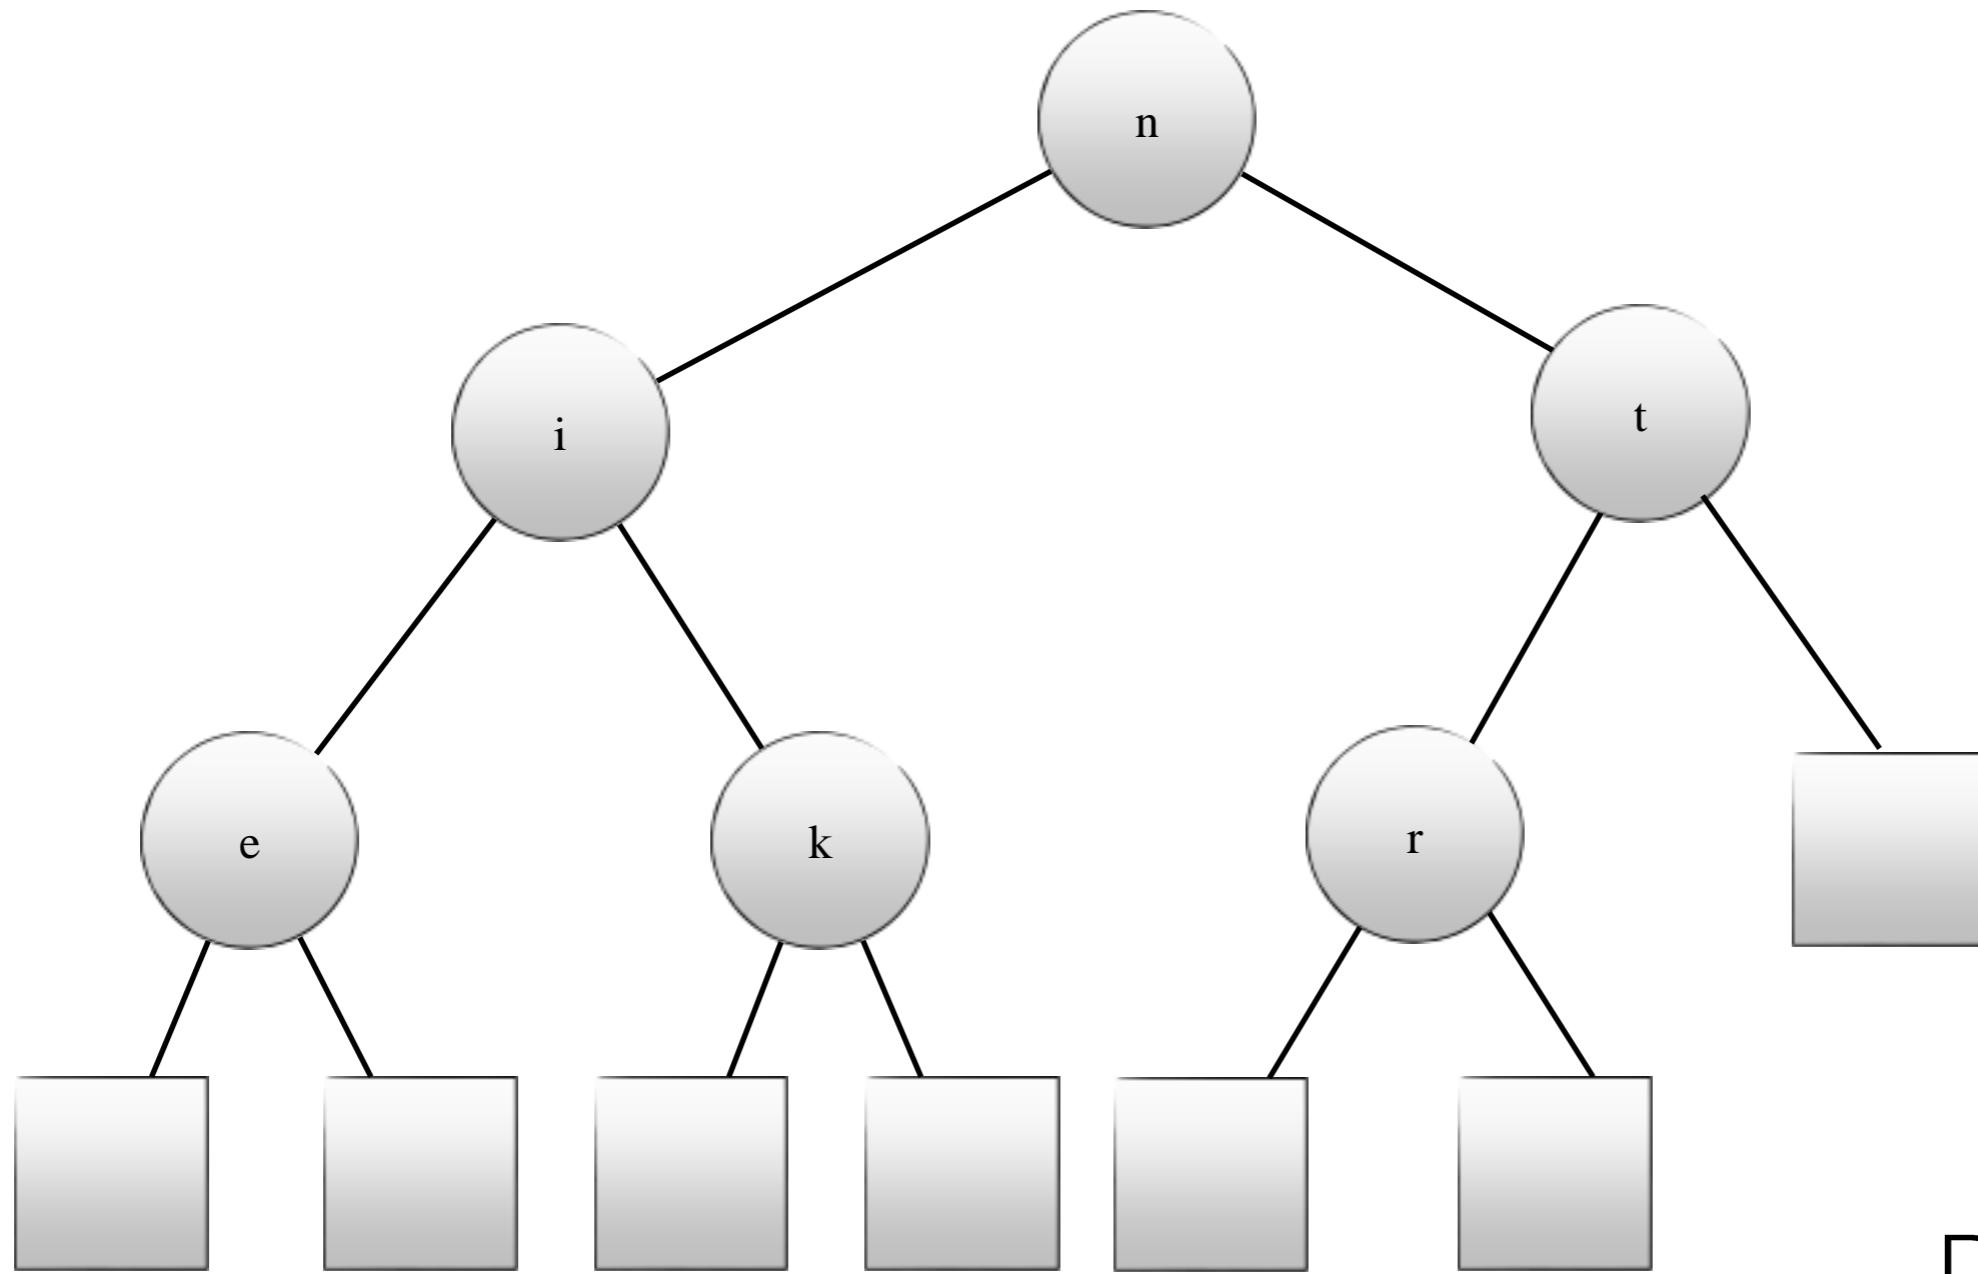
# Delete key s



# Delete key s

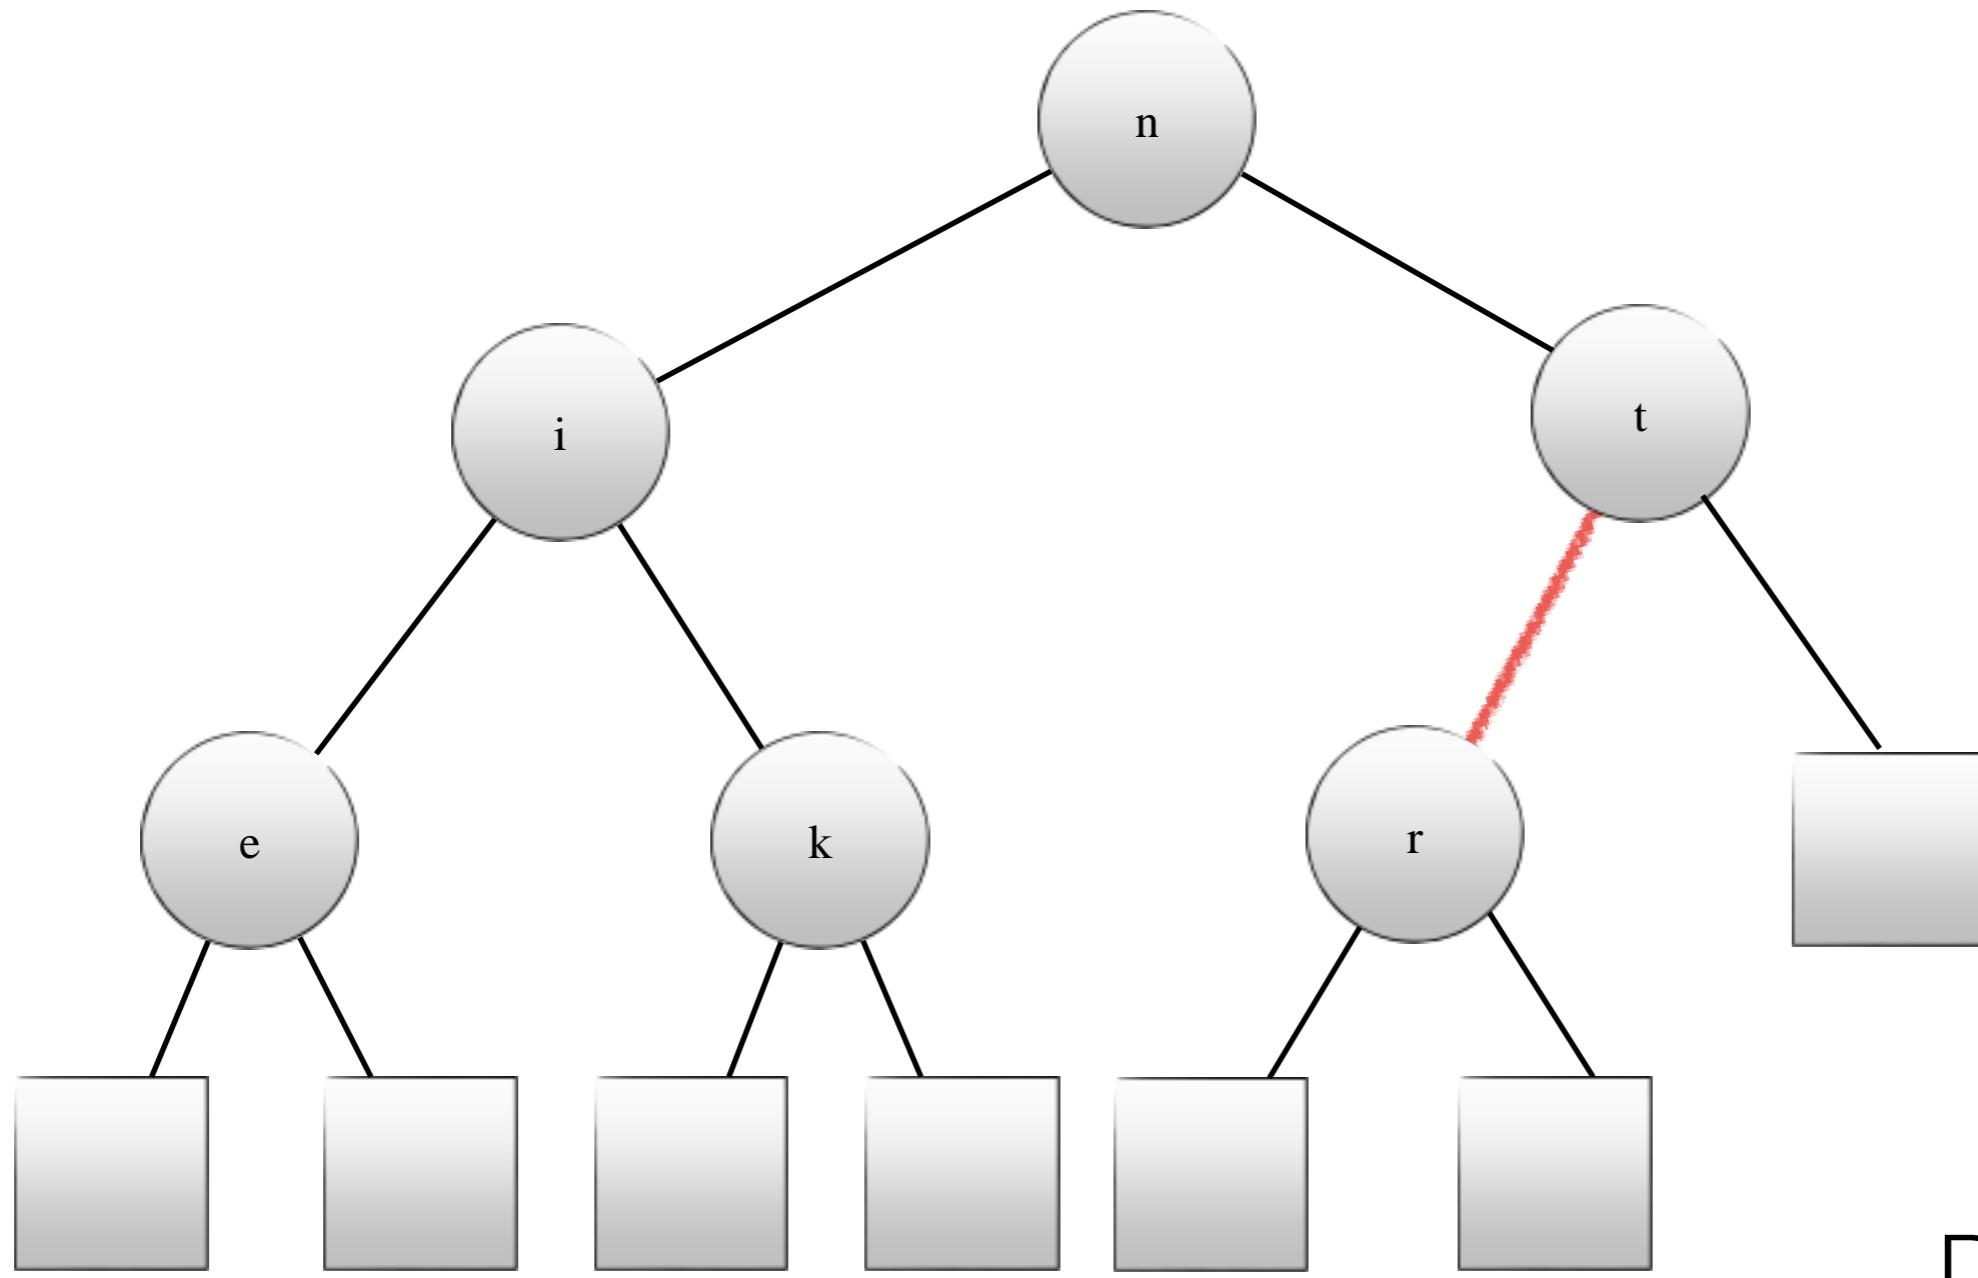


# Delete key s



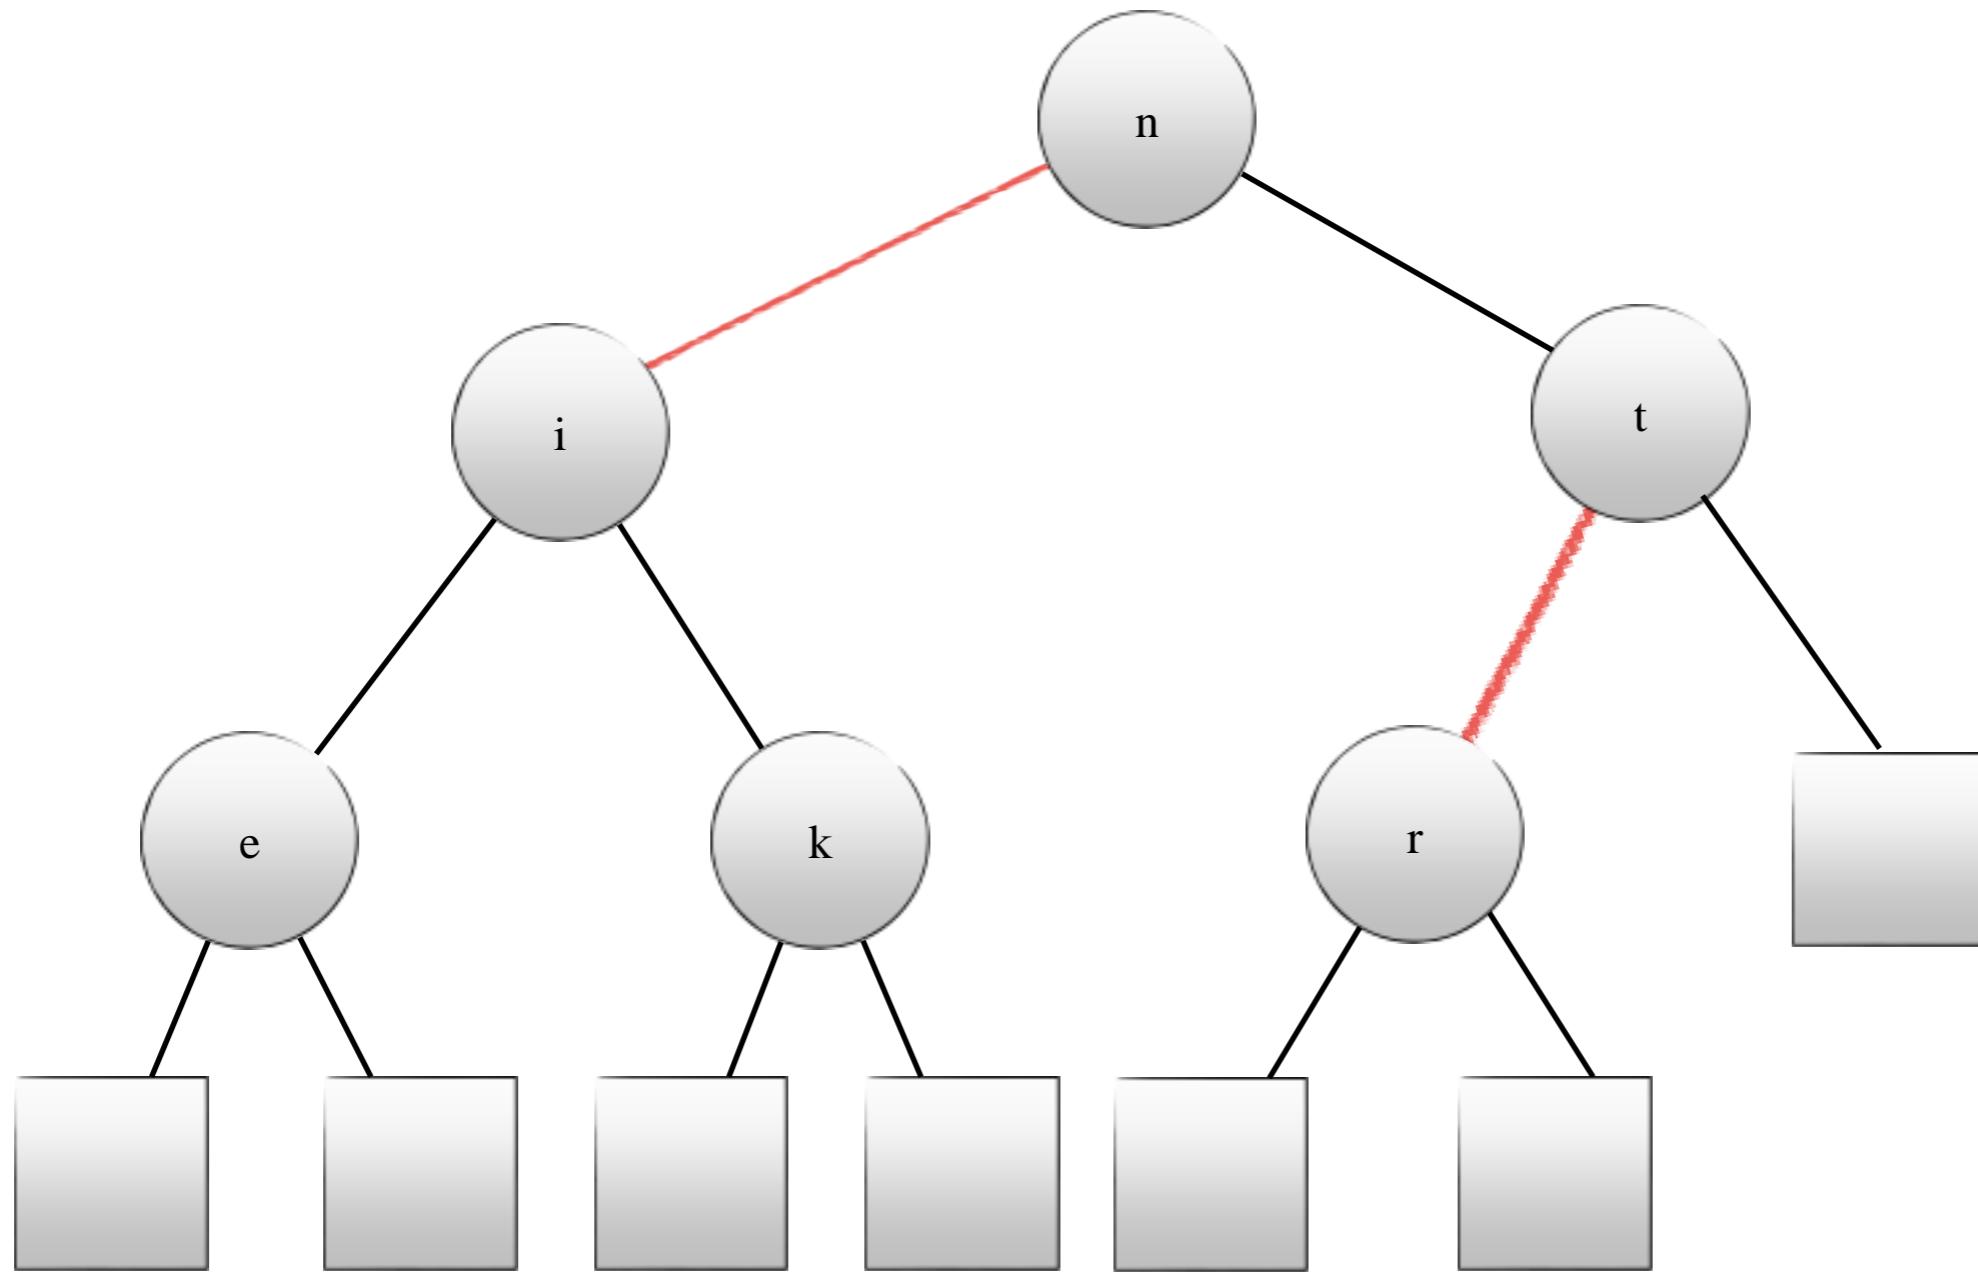
Doesn't work

# Delete key s



Doesn't work

# Delete key s



We need!

# Deletion from 2-3/red-black trees

- Deletion is slightly more complex than insertion as you need to account for a potential imbalance on the way “down”
- So, we rotate and/or color flip both on the way down the recursion as well as on the way up.
- Generally, we make 2-nodes into 3-nodes or 4-nodes on the way down to accommodate a removal then correct the red edges on the way back up if needed.

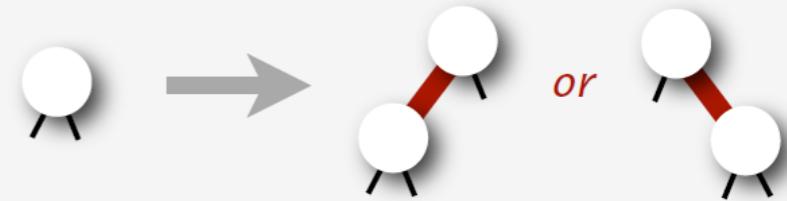
# Insert code for LLRB trees

Introduction  
2-3-4 Trees  
Red-Black Trees  
Left-Leaning RB Trees  
Deletion

is based on three simple operations.

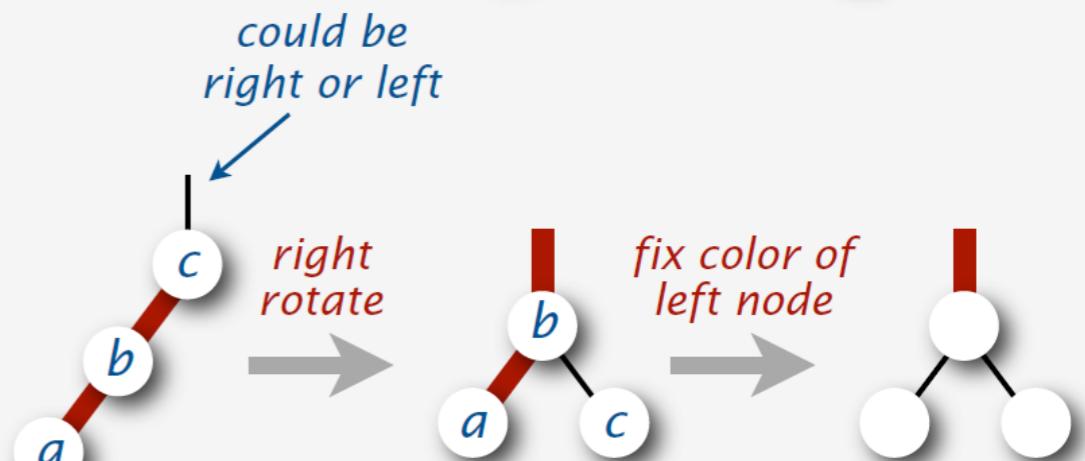
1. Insert a new node at the bottom.

```
if (h == null)
    return new Node(key, value, RED);
```



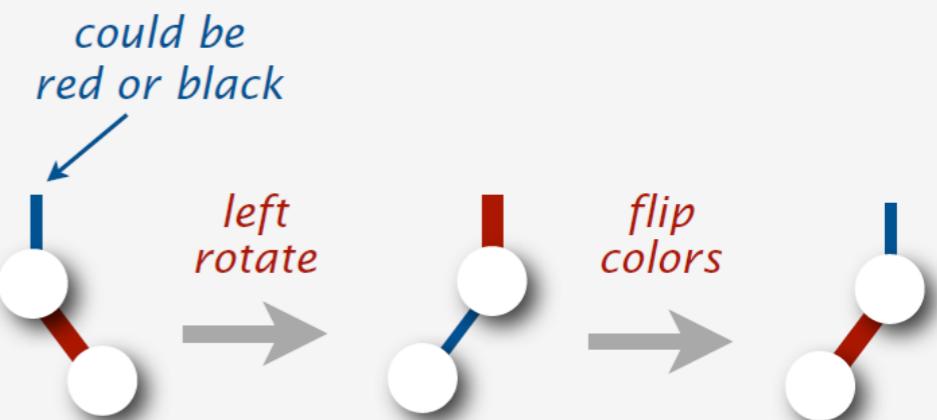
2. Split a 4-node.

```
private Node splitFourNode(Node h)
{
    x = rotR(h);
    x.left.color = BLACK;
    return x;
}
```



3. Enforce left-leaning condition.

```
private Node leanLeft(Node h)
{
    x = rotL(h);
    x.color      = x.left.color;
    x.left.color = RED;
    return x;
}
```



# Insert implementation for LLRB trees

is a few lines of code added to elementary BST insert

Introduction  
2-3-4 Trees  
Red-Black Trees  
Left-Leaning RB Trees  
Deletion

```
private Node insert(Node h, Key key, Value val)
{
    if (h == null)
        return new Node(key, val, RED);           ← insert at the bottom

    if (isRed(h.left))
        if (isRed(h.left.left))
            h = splitFourNode(h);                ← split 4-nodes on the way down

    int cmp = key.compareTo(h.key);
    if (cmp == 0) h.val = val;
    else if (cmp < 0)
        h.left = insert(h.left, key, val);       ← standard BST insert code
    else
        h.right = insert(h.right, key, val);

    if (isRed(h.right))
        h = leanLeft(h);                      ← fix right-leaning reds on the way up

    return h;
}
```

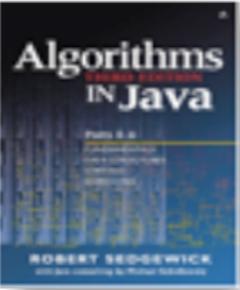
# Why revisit red-black trees?

Introduction  
2-3-4 Trees  
Red-Black Trees  
Left-Leaning RB Trees  
Deletion

Take your pick:

```
private Node insert(Node x, Key key, Value val, boolean sw)
{
    if (x == null)
        return new Node(key, value, RED);
    int cmp = key.compareTo(x.key);

    if (isRed(x.left) && isRed(x.right))
    {
        x.color = RED;
        x.left.color = BLACK;
        x.right.color = BLACK;
    }
    if (cmp == 0) x.val = val;
    else if (cmp < 0)
    {
        x.left = insert(x.left, key, val, false);
        if (isRed(x) && isRed(x.left) && sw)
            x = rotR(x);
        if (isRed(x.left) && isRed(x.left.left))
        {
            x = rotR(x);
            x.color = BLACK; x.right.color = RED;
        }
    }
    else // if (cmp > 0)
    {
        x.right = insert(x.right, key, val, true);
        if (isRed(h) && isRed(x.right) && !sw)
            x = rotL(x);
        if (isRed(h.right) && isRed(h.right.right))
        {
            x = rotL(x);
            x.color = BLACK; x.left.color = RED;
        }
    }
    return x;
}
```



```
private Node insert(Node h, Key key, Value val)
{
    int cmp = key.compareTo(h.key);
    if (h == null)
        return new Node(key, val, RED);
    if (isRed(h.left))
        if (isRed(h.left.left))
        {
            h = rotR(h);
            h.left.color = BLACK;
        }
    if (cmp == 0) x.val = val;
    else if (cmp < 0)
        h.left = insert(h.left, key, val);
    else
        h.right = insert(h.right, key, val);
    if (isRed(h.right))
    {
        h = rotL(h);
        h.color = h.left.color;
        h.left.color = RED;
    }
    return h;
}
```

**Left-Leaning  
Red-Black Trees**  
Robert Sedgewick  
Princeton University

straightforward

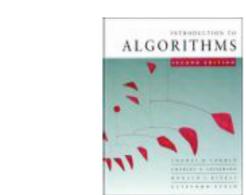
very  
tricky

# Why revisit red-black trees?

Take your pick:

TreeMap.java

Adapted from  
CLR by  
experienced  
professional  
programmers  
(2004)



150

wrong scale!

Why left-leaning trees?

Take your pick:

```
private Node insert(Node x, Key key, Value val, boolean sw)
{
    if (x == null)
        return new Node(key, value, RED);
    int cmp = key.compareTo(x.key);
    if (isRed(x.left) && isRed(x.right))
    {
        x.color = RED;
        x.left.color = BLACK;
        x.right.color = BLACK;
    }
    if (cmp == 0) x.val = val;
    else if (cmp < 0)
    {
        x.left = insert(x.left, key, val, false);
        if (isRed(x) && isRed(x.left) && sw)
            x = rotR(x);
        if (isRed(x.left) && isRed(x.left.left))
        {
            x = rotR(x);
            x.color = BLACK; x.right.color = RED;
        }
    } else // if (cmp > 0)
    {
        x.right = insert(x.right, key, val, true);
        if (isRed(h) && isRed(x.right) && !sw)
            x = rotL(x);
        if (isRed(h.right) && isRed(h.right.right))
        {
            x = rotL(x);
            x.color = BLACK; x.left.color = RED;
        }
    }
    return x;
}
```



```
private Node insert(Node h, Key key, Value val)
{
    int cmp = key.compareTo(h.key);
    if (h == null)
        return new Node(key, val, RED);
    if (isRed(h.left))
        if (isRed(h.left.left))
        {
            h = rotR(h);
            h.left.color = BLACK;
        }
    if (cmp == 0) h.val = val;
    else if (cmp < 0)
        h.left = insert(h.left, key, val);
    else
        h.right = insert(h.right, key, val);
    if (isRed(h.right))
    {
        h = rotL(h);
        h.color = h.left.color;
        h.left.color = RED;
    }
    return h;
}
```

Introduction  
2-3-4 Trees  
Red-Black Trees  
Left-Leaning RB Trees  
Deletion

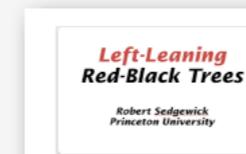
Left-Leaning  
Red-Black Trees  
Robert Sedgewick  
Princeton University

straightforward

very  
tricky



40



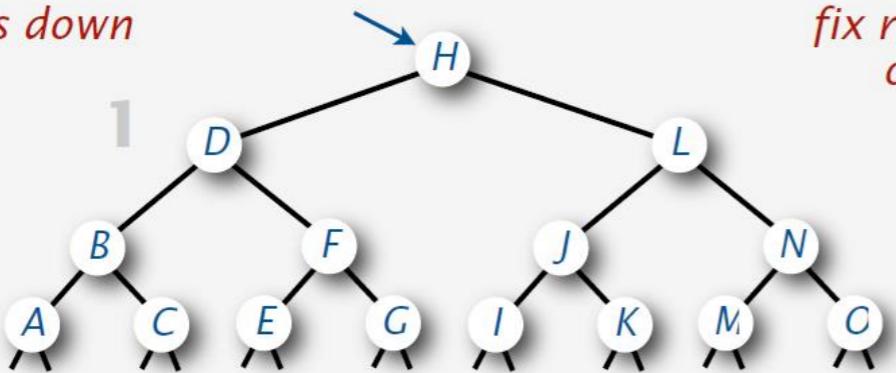
30

← lines of code for insert  
(lower is better!)

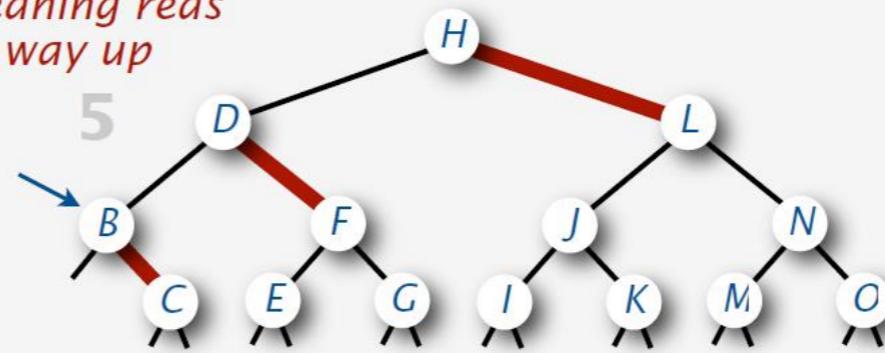
# deleteMin() example

Introduction  
2-3-4 Trees  
Red-Black Trees  
Left-Leaning RB Trees  
Deletion

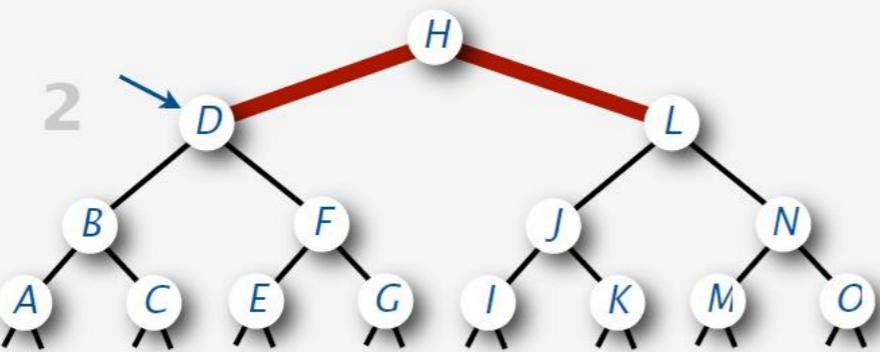
push reds down



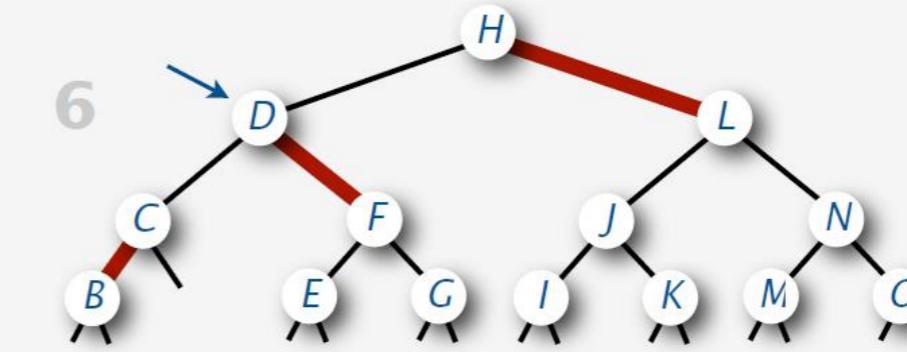
fix right-leaning reds  
on the way up



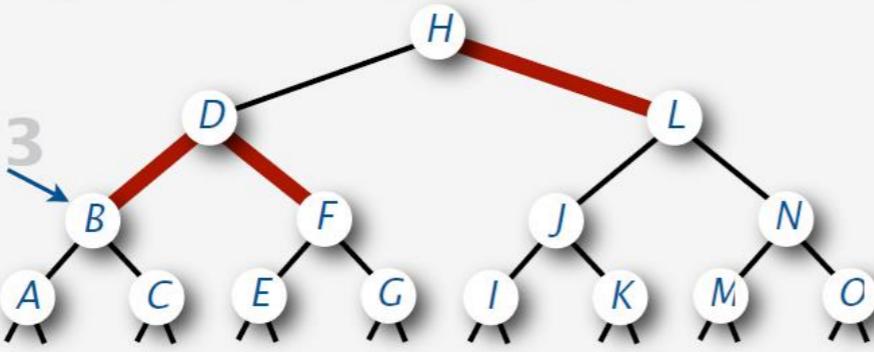
2



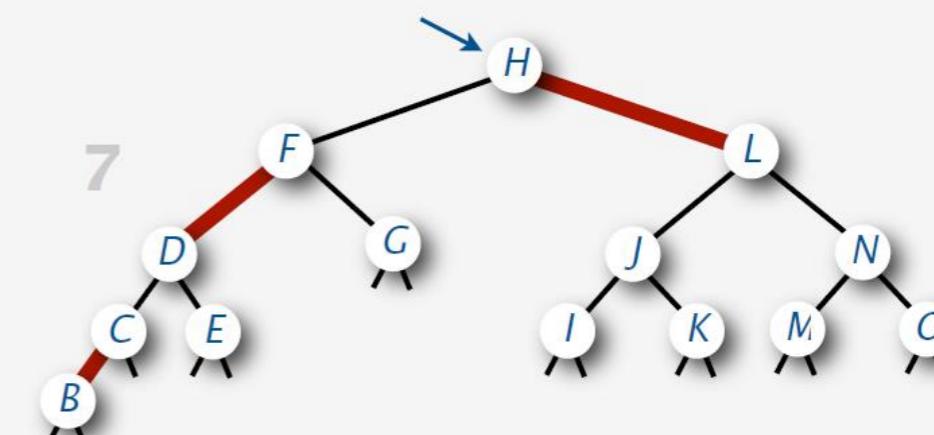
6



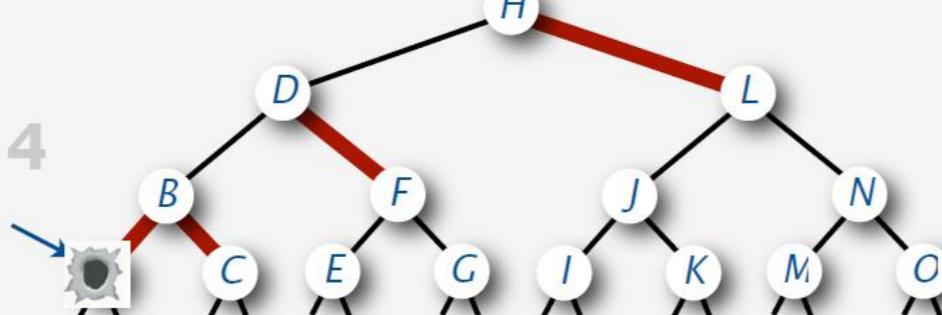
3



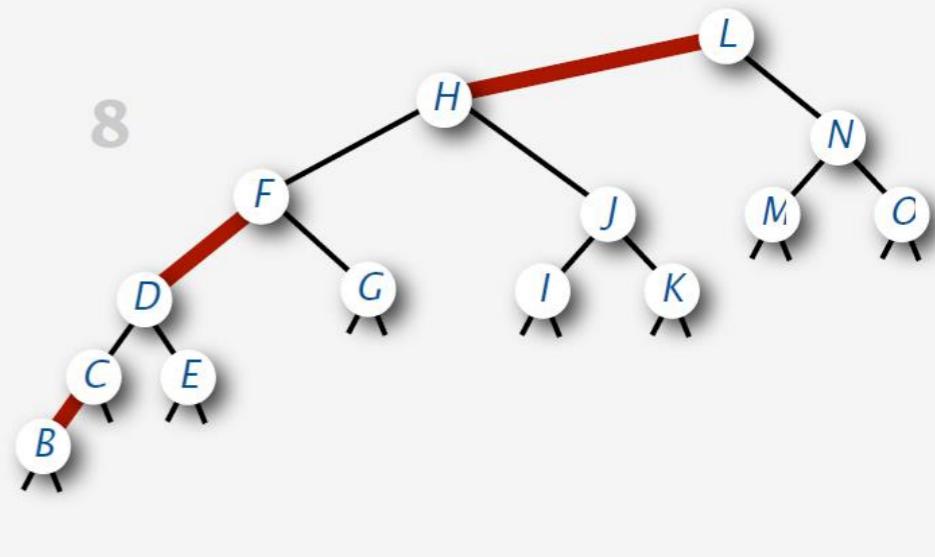
7



4

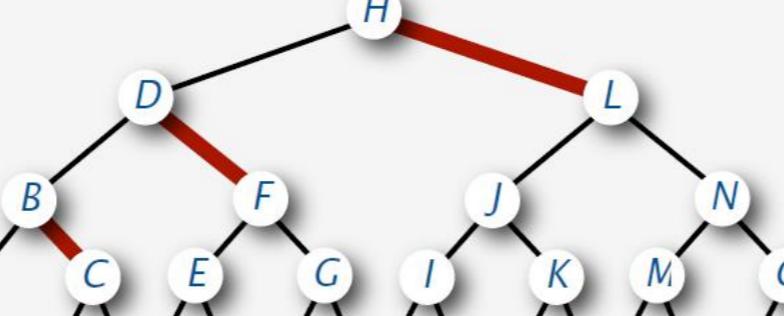


8



remove minimum

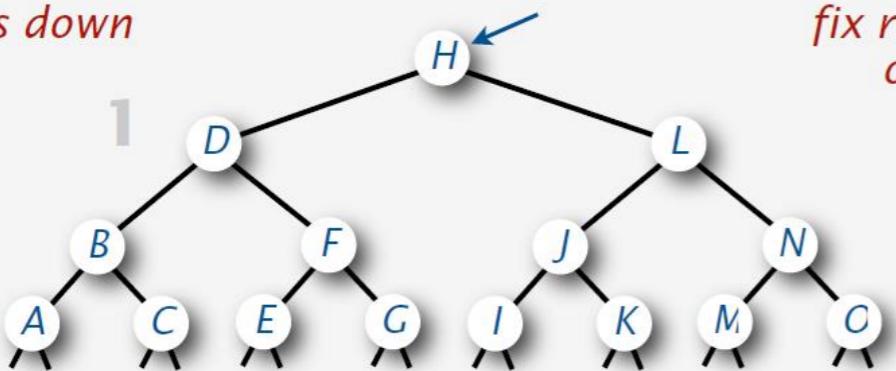
5



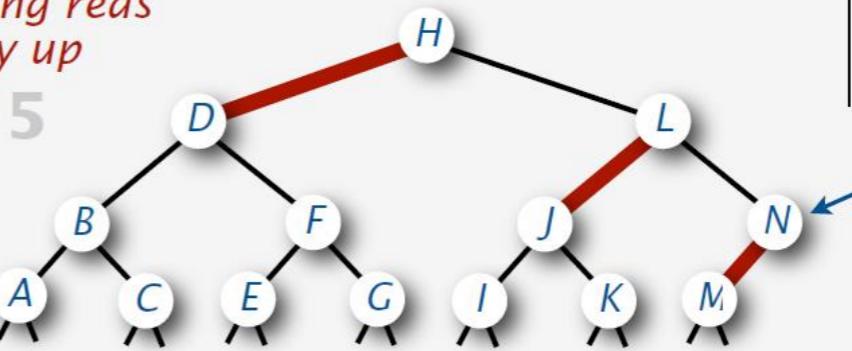
# deleteMax() example

Introduction  
2-3-4 Trees  
Red-Black Trees  
Left-Leaning RB Trees  
Deletion

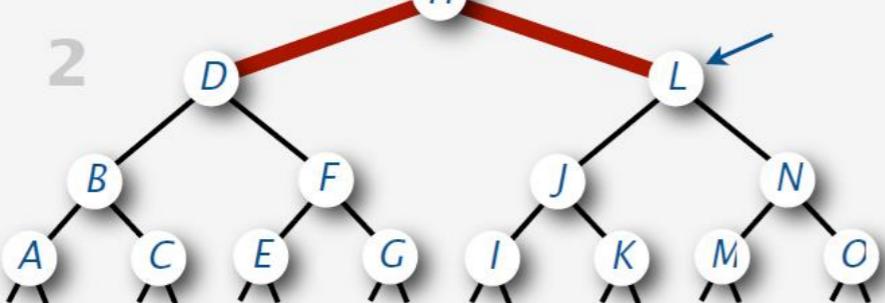
push reds down



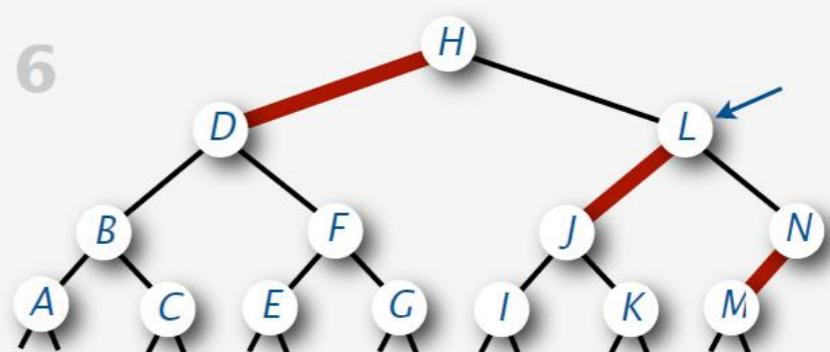
fix right-leaning reds  
on the way up



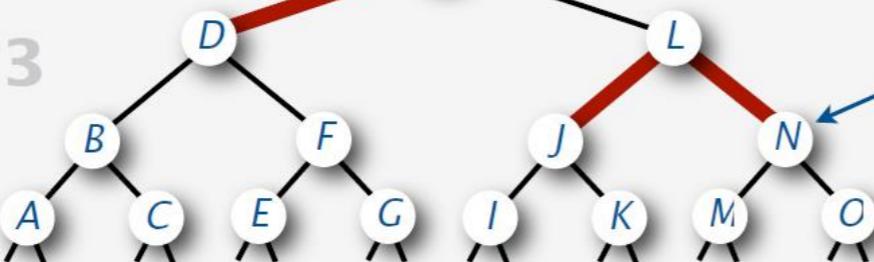
2



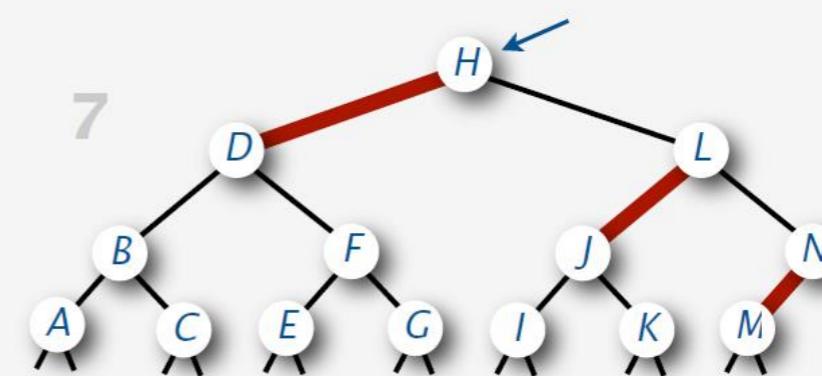
6



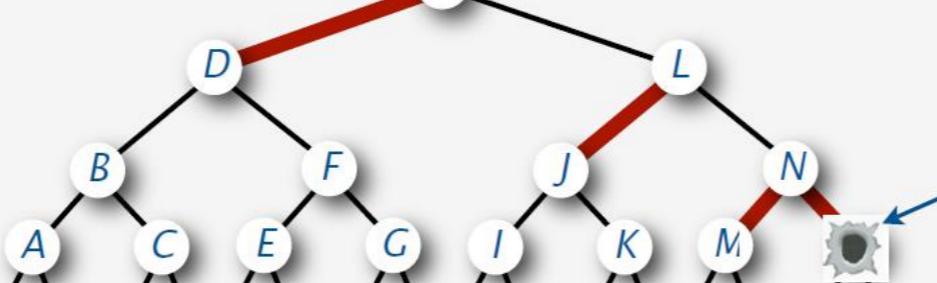
3



7

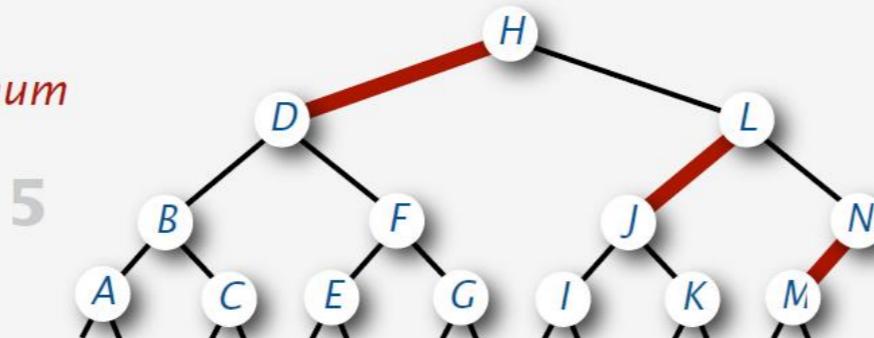


4



(nothing to fix!)

remove maximum

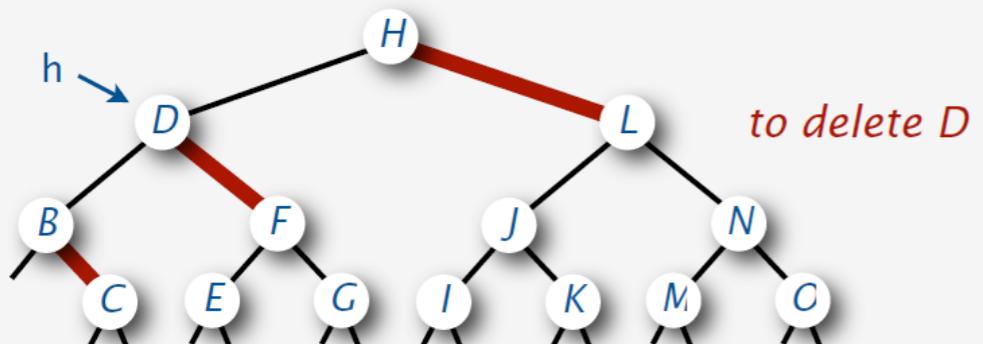


# Deleting an arbitrary node

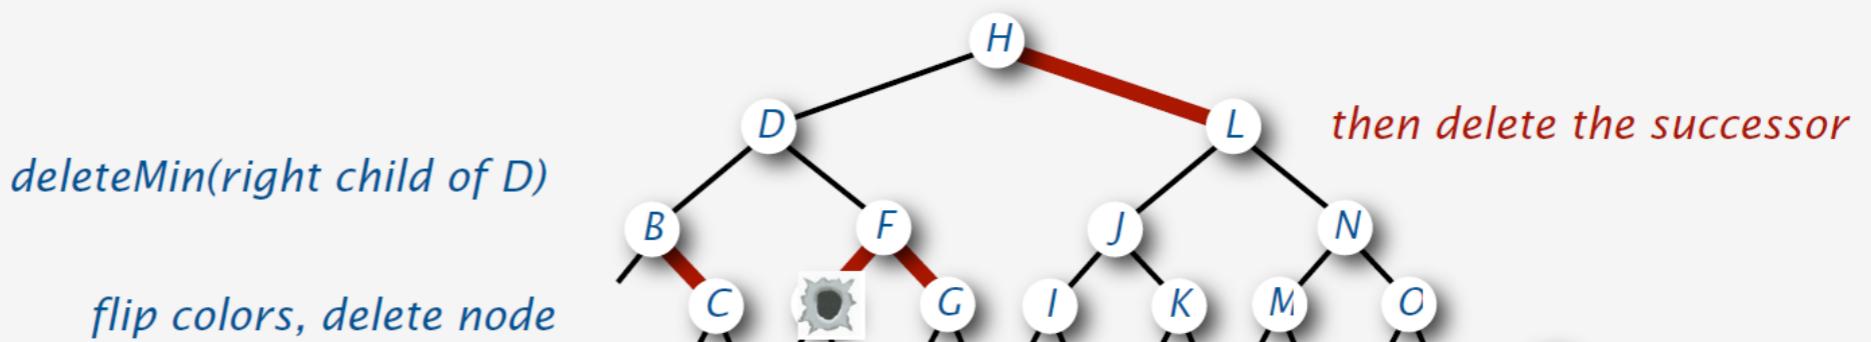
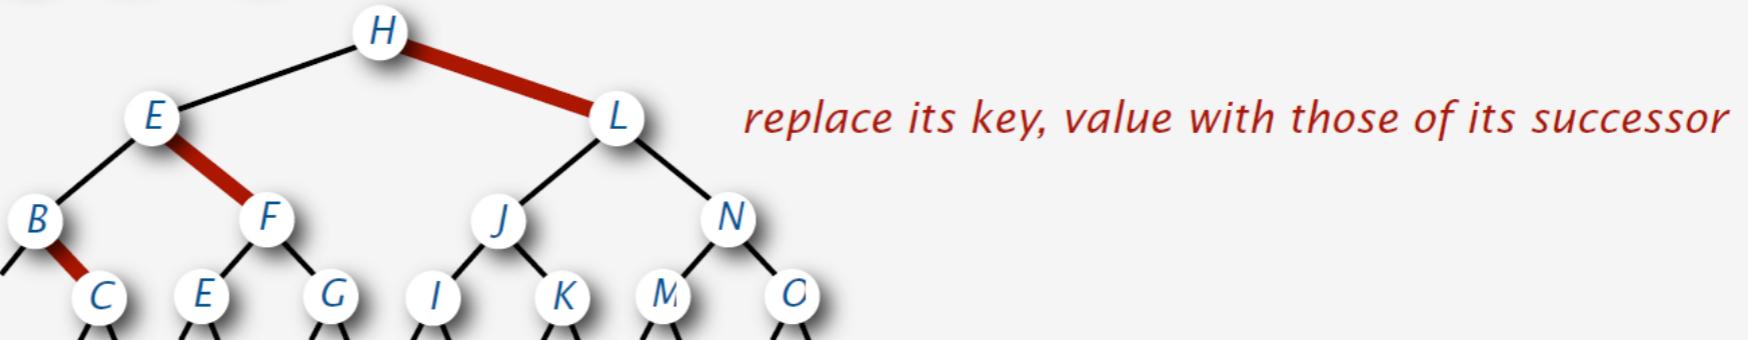
reduces to `deleteMin()`

Introduction  
2-3-4 Trees  
Red-Black Trees  
Left-Leaning RB Trees  
Deletion

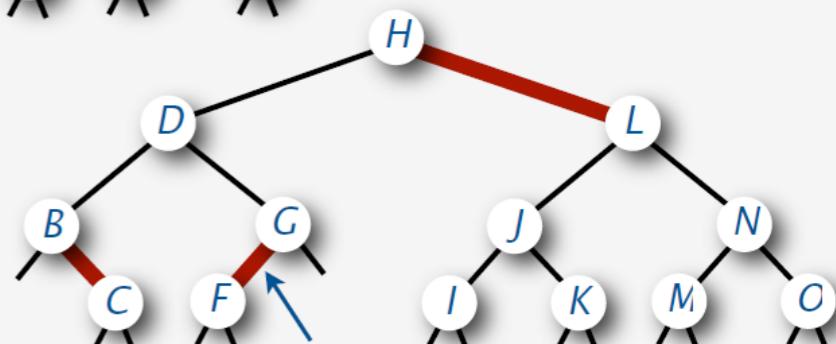
A standard trick:



```
h.key    = min(h.right);  
h.value = get(h.right, h.key);  
h.right = deleteMin(h.right);
```



*fix right-leaning red link*



# Need MORE Demos?

- [http://inst.eecs.berkeley.edu/~cs61b/fa17/materials/  
demos/ll-red-black-demo.html](http://inst.eecs.berkeley.edu/~cs61b/fa17/materials/demos/ll-red-black-demo.html)

# Red-Black trees in the wild

---

Red-black trees are widely used as system symbol tables.

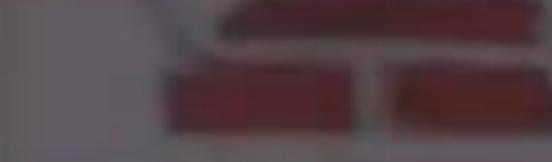
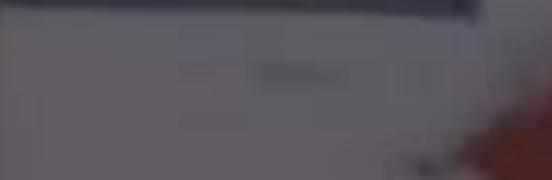
- Java: `java.util.TreeMap`, `java.util.TreeSet`.
- C++ STL: `map`, `multimap`, `multiset`.
- Linux kernel: completely fair scheduler, `linux/rbtree.h`.
- Emacs: conservative stack scanning.



*Common sense. Sixth sense.*

*Together they're the  
FBI's newest team.*

VISITOR  
MAP



# Red-black BSTs in the wild

---

## ACT FOUR

FADE IN:

48 INT. FBI HQ - NIGHT

48

Antonio is at THE COMPUTER as Jess explains herself to Nicole and Pollock. The CONFERENCE TABLE is covered with OPEN REFERENCE BOOKS, TOURIST GUIDES, MAPS and REAMS OF PRINTOUTS.

JESS

It was the red door again.

POLLOCK

I thought the red door was the storage container.

JESS

But it wasn't red anymore. It was black.

ANTONIO

So red turning to black means... what?

POLLOCK

Budget deficits? Red ink, black ink?

NICOLE

Yes. I'm sure that's what it is. But maybe we should come up with a couple other options, just in case.

Antonio refers to his COMPUTER SCREEN, which is filled with mathematical equations.

ANTONIO

It could be an algorithm from a binary search tree. A red-black tree tracks every simple path from a node to a descendant leaf with the same number of black nodes.

JESS

Does that help you with girls?

## B-trees (Bayer-McCreight, 1972)

B-Tree. Generalizes 2-3-4 trees by allowing up to  $M$  links per node.

Main application: file systems.

- Reading a page into memory from disk is expensive.
- Accessing info on a page in memory is free.
- Goal: minimize # page accesses.
- Node size  $M$  = page size.

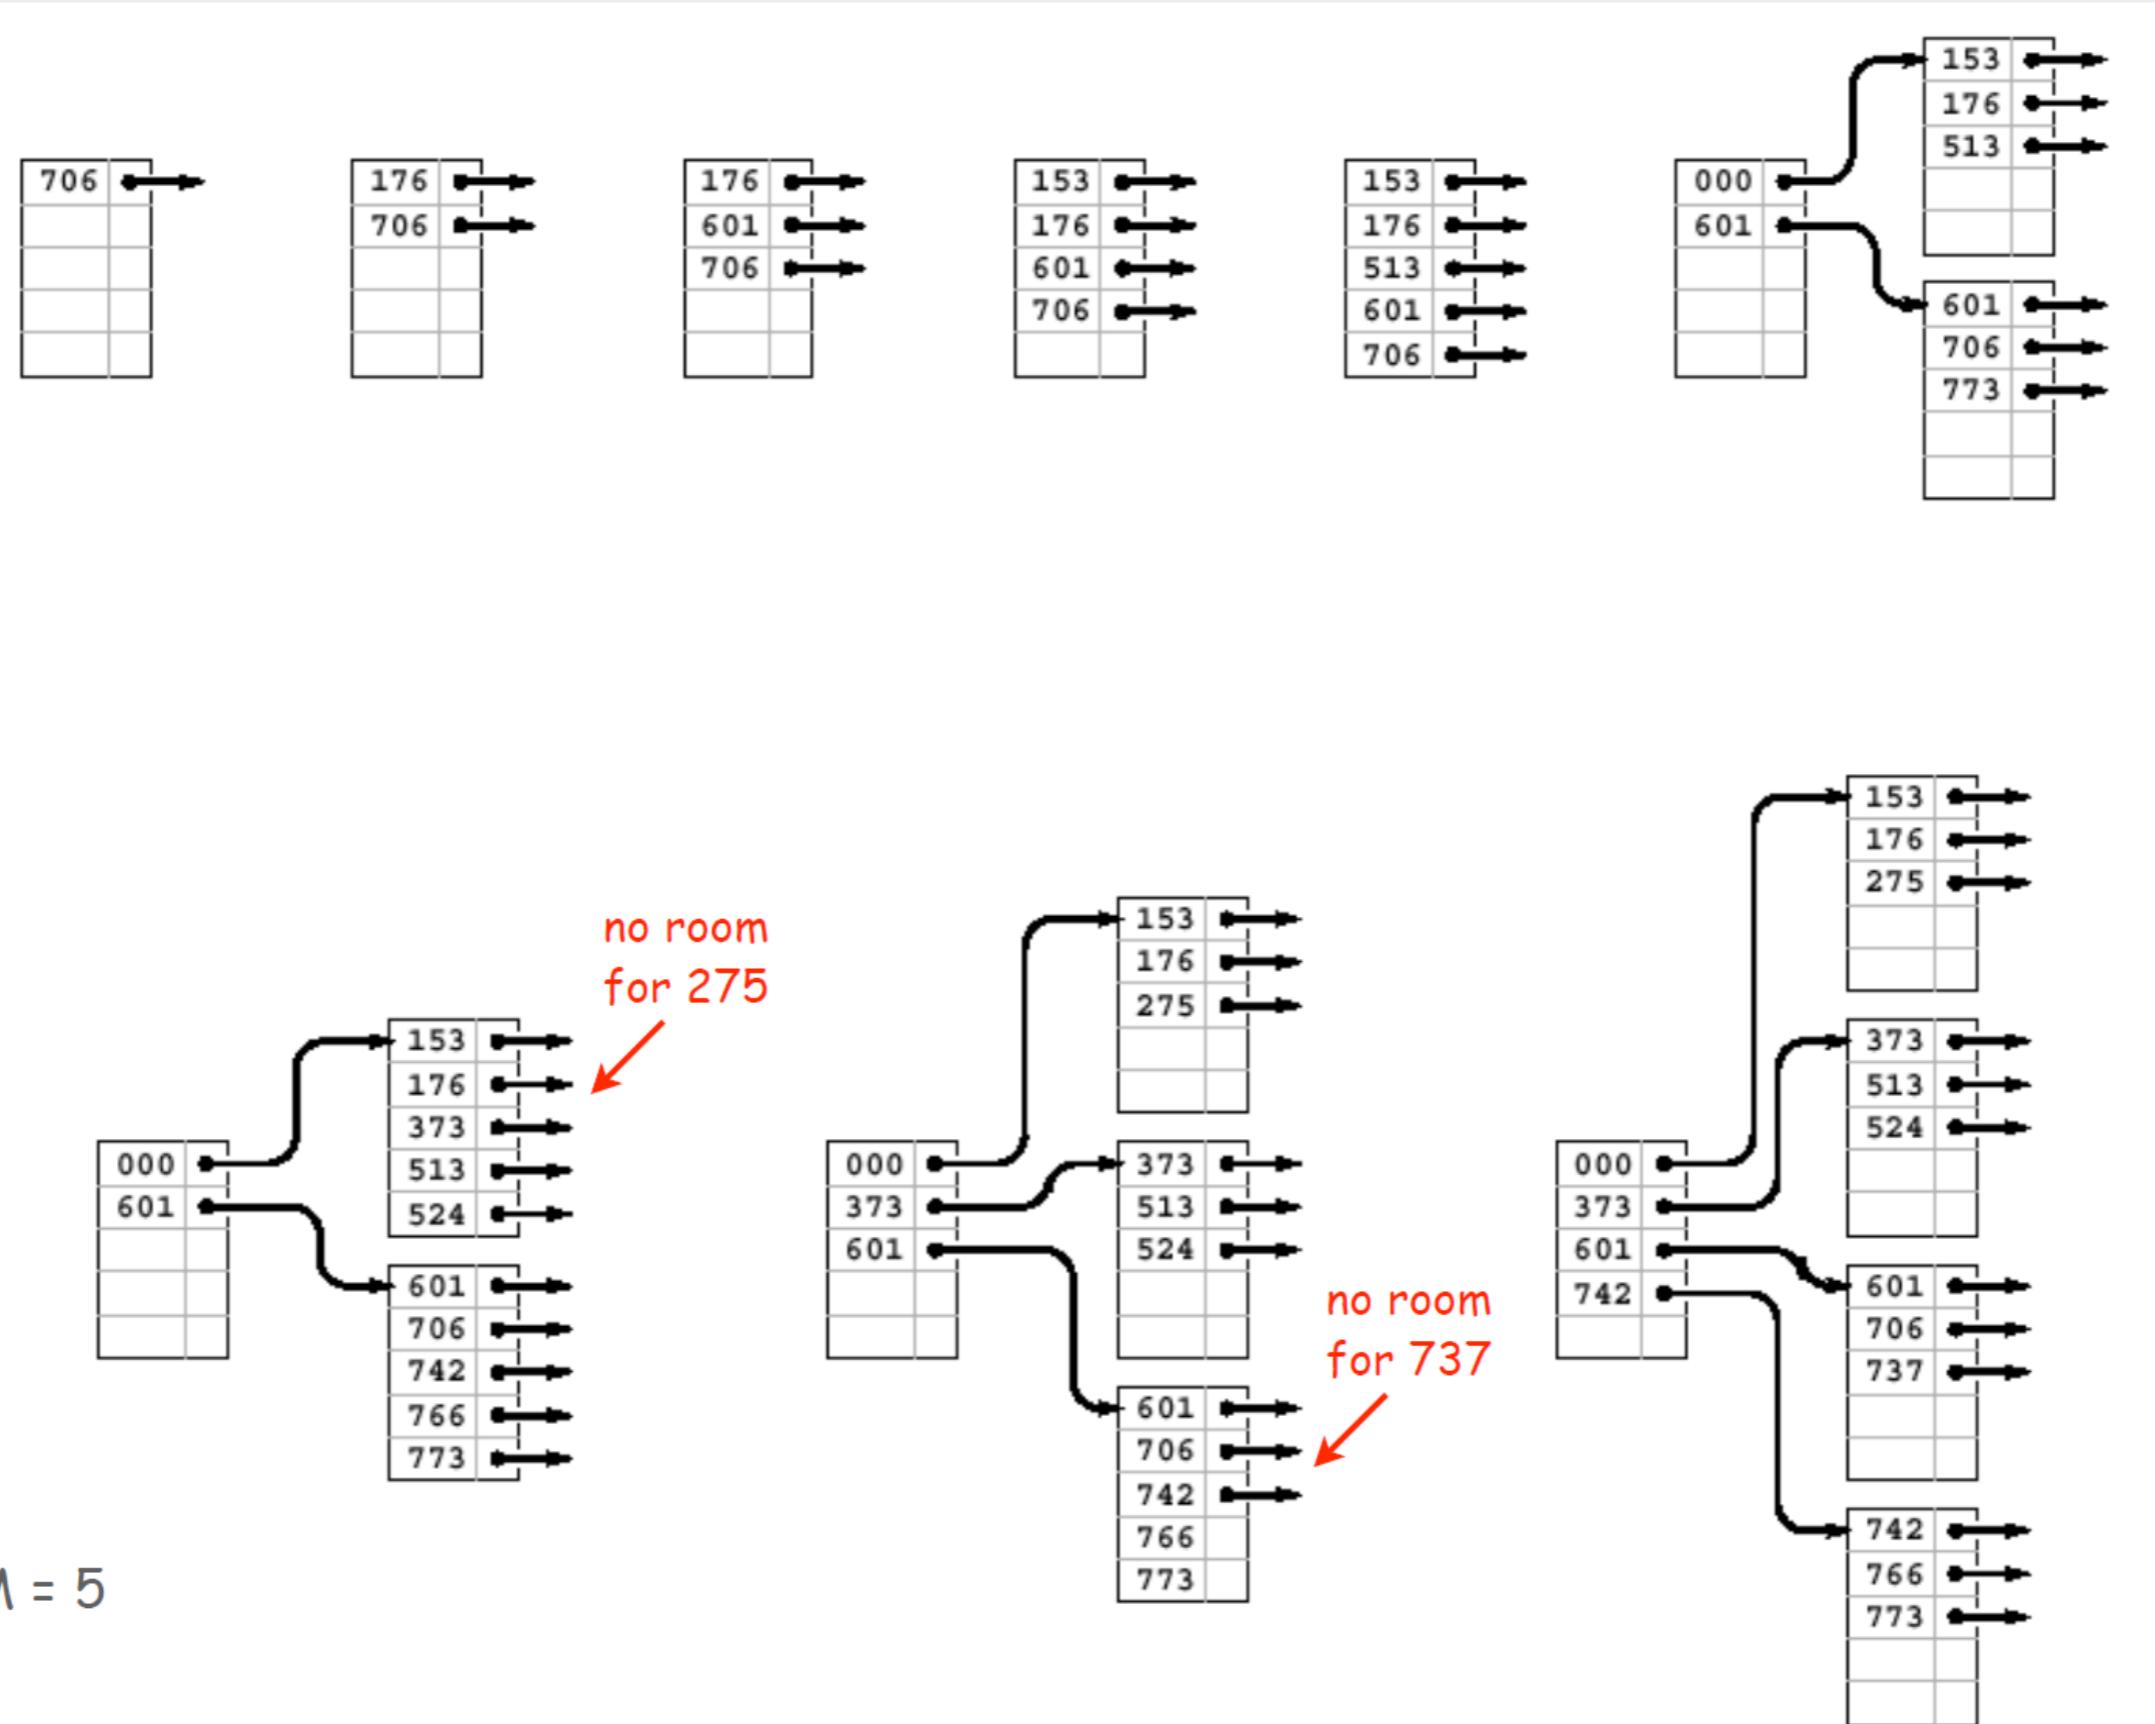
Space-time tradeoff.

- $M$  large  $\Rightarrow$  only a few levels in tree.
- $M$  small  $\Rightarrow$  less wasted space.
- Typical  $M = 1000$ ,  $N < 1$  trillion.

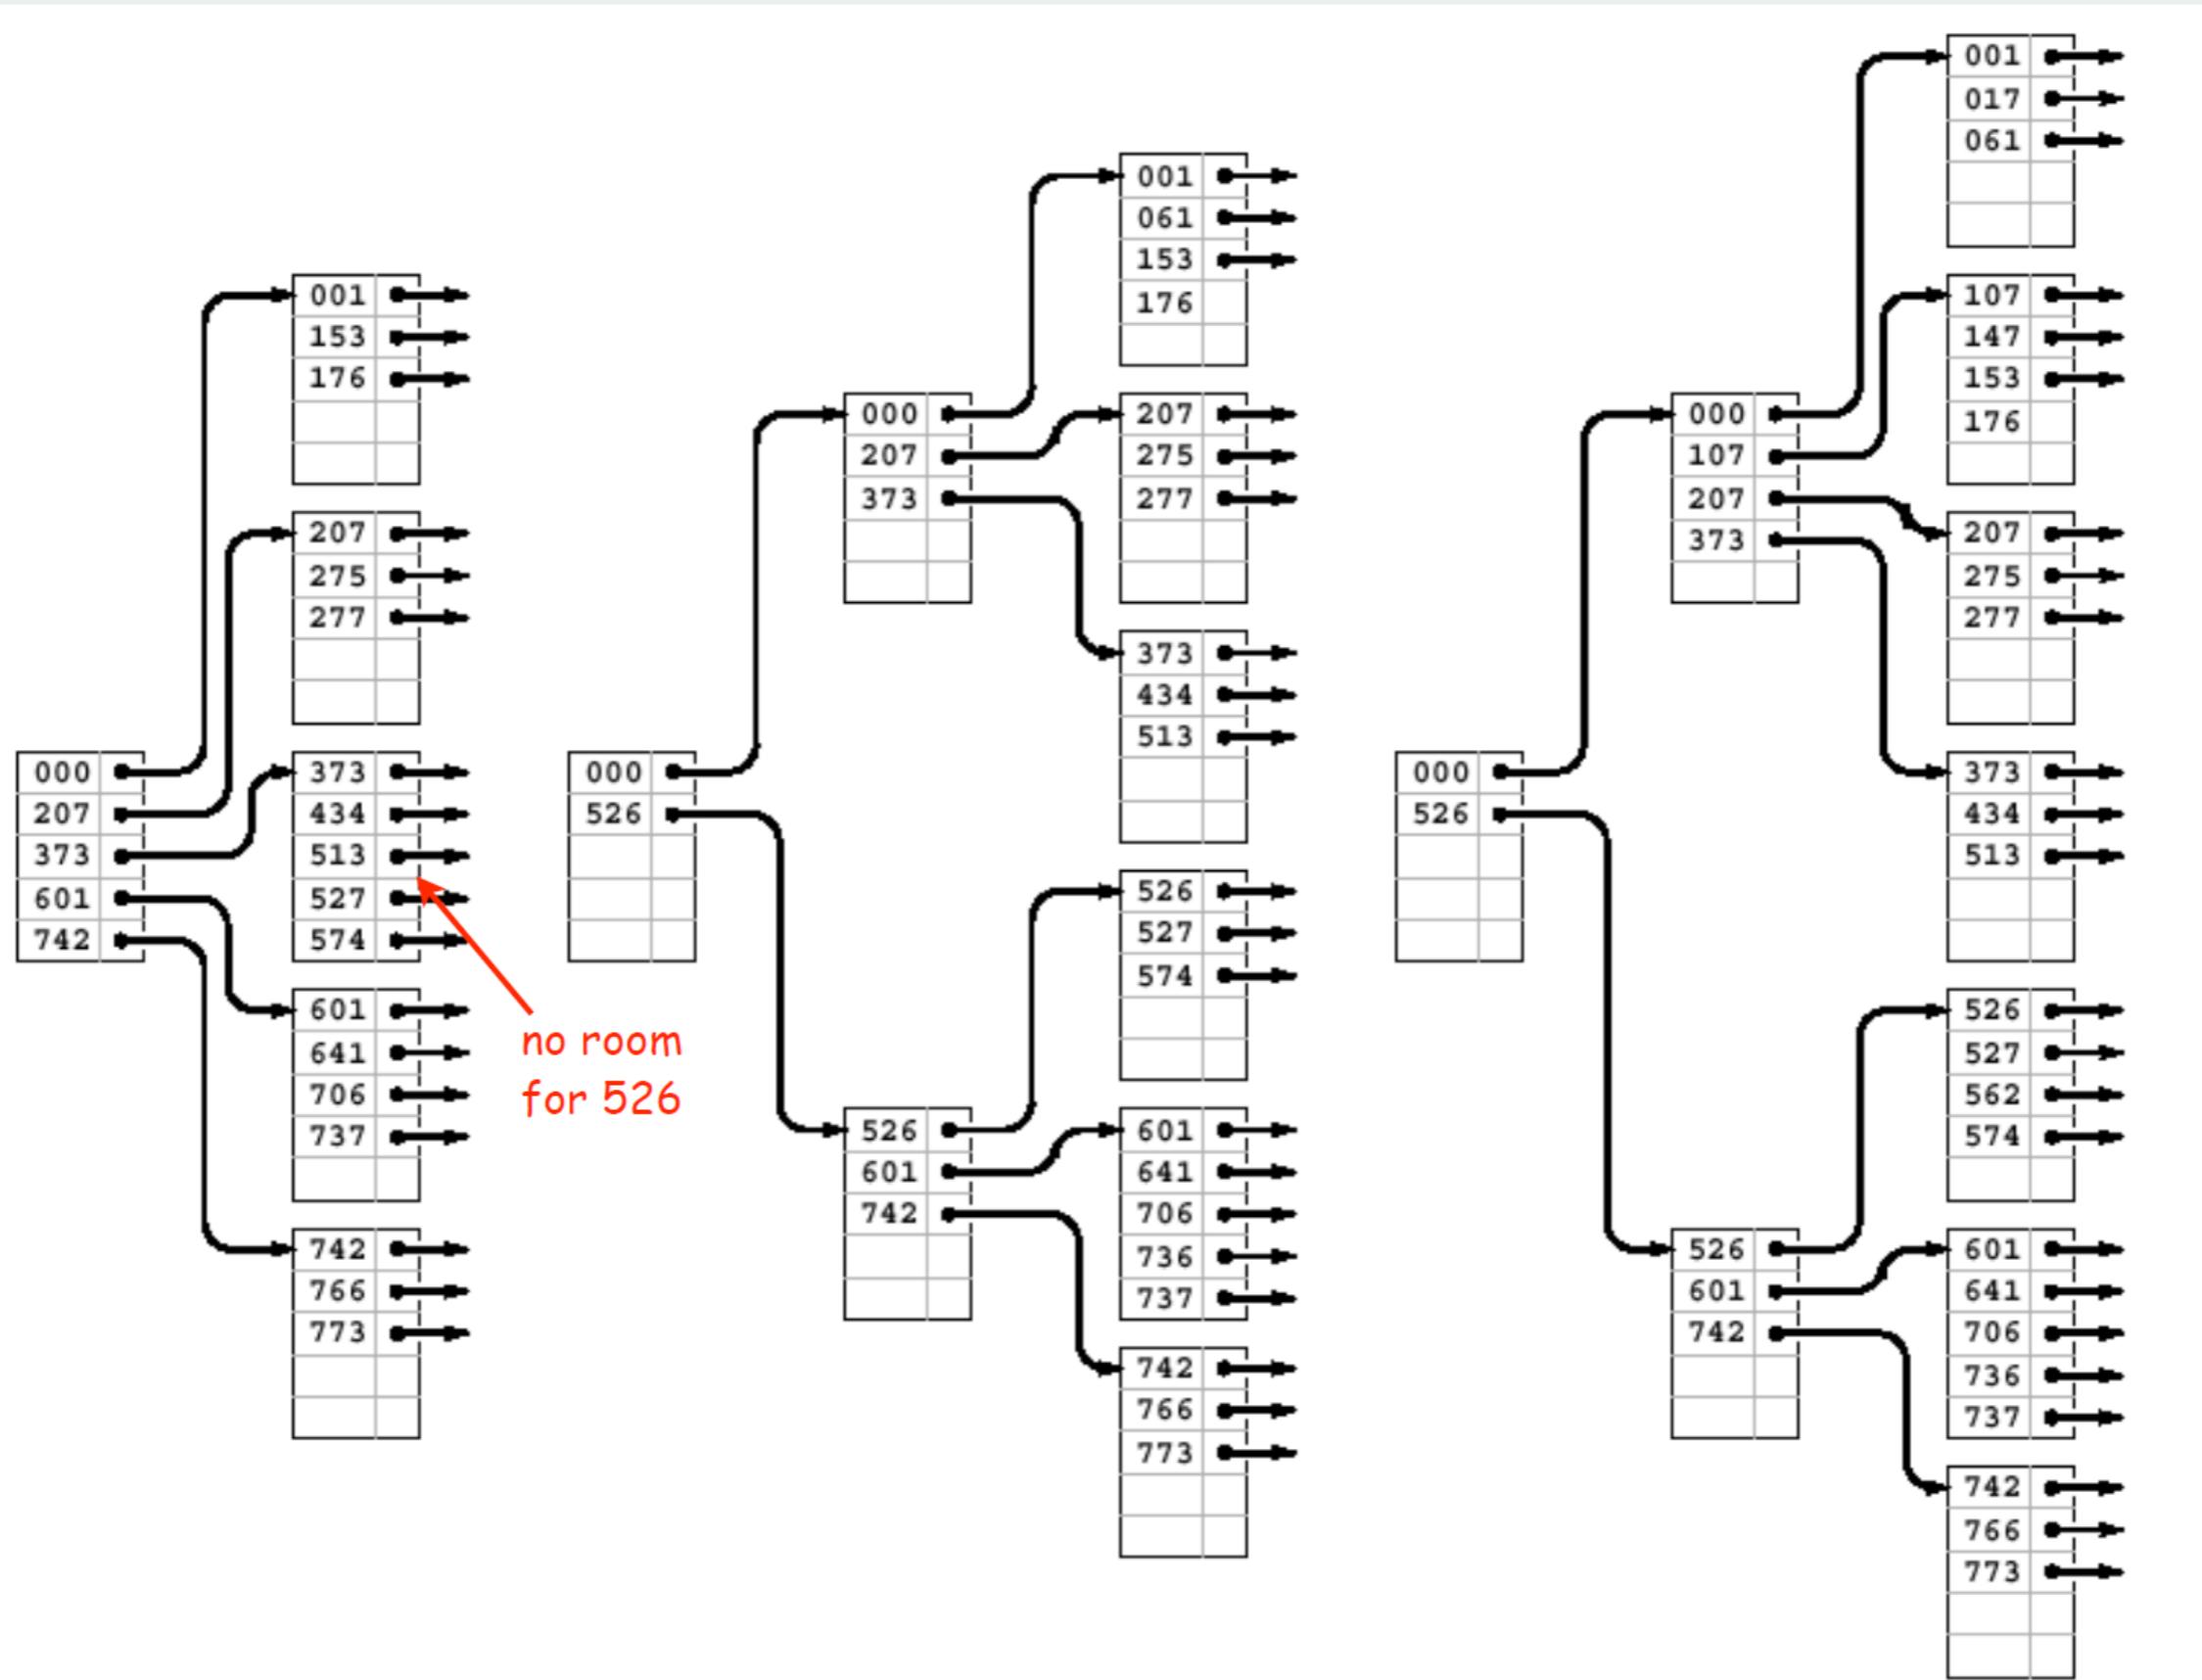
Bottom line. Number of **page** accesses is  $\log_M N$  per op.

↑  
in practice: 3 or 4 (!)

## B-Tree Example



## B-Tree Example (cont)



## Summary of symbol-table implementations

implementation	guarantee			average case			ordered iteration?
	search	insert	delete	search	insert	delete	
unordered array	N	N	N	N/2	N/2	N/2	no
ordered array	$\lg N$	N	N	$\lg N$	N/2	N/2	yes
unordered list	N	N	N	N/2	N	N/2	no
ordered list	N	N	N	N/2	N/2	N/2	yes
BST	N	N	N	$1.44 \lg N$	$1.44 \lg N$	?	yes
randomized BST	$7 \lg N$	$7 \lg N$	$7 \lg N$	$1.44 \lg N$	$1.44 \lg N$	$1.44 \lg N$	yes
2-3-4 tree	$c \lg N$	$c \lg N$		$c \lg N$	$c \lg N$		yes
red-black tree	$2 \lg N$	$2 \lg N$	$2 \lg N$	$\lg N$	$\lg N$	$\lg N$	yes
B-tree	1	1	1	1	1	1	yes

B-Tree. Number of page accesses is  $\log_M N$  per op.

## Balanced trees in the wild

**Red-black trees:** widely used as system symbol tables

- Java: `java.util.TreeMap`, `java.util.TreeSet`.
- C++ STL: `map`, `multimap`, `multiset`.
- Linux kernel: `linux/rbtree.h`.

**B-Trees:** widely used for file systems and databases

- Windows: HPFS.
- Mac: HFS, HFS+.
- Linux: ReiserFS, XFS, Ext3FS, JFS.
- Databases: ORACLE, DB2, INGRES, SQL, PostgreSQL

**Bottom line:** ST implementation with  $\lg N$  **guarantee** for all ops.

- Algorithms are variations on a theme: rotations when inserting.
- Easiest to implement, optimal, fastest in practice: **LLRB trees**
- Abstraction extends to give search algorithms for huge files: **B-trees**